



# ОБРАБОТКА ОЗЕРА ДАННЫХ НА NODE.JS В SERVERLESS-АРХИТЕКТУРЕ



**Николай Матвиенко**

**Grid Dynamics**






# Nikolay Matvienko

Senior Software Engineer and Node.js expert at Grid Dynamics  
Node.js diagnostics, performance improvement consultant.

You can find me at [twitter.com/matvi3nko](https://twitter.com/matvi3nko)  
[github.com/matvi3nko](https://github.com/matvi3nko)



# DATA LAKE




RAW DATA MODEL

STRUCTURED DATA

ANALYZED DATA MODEL



# THE GOAL



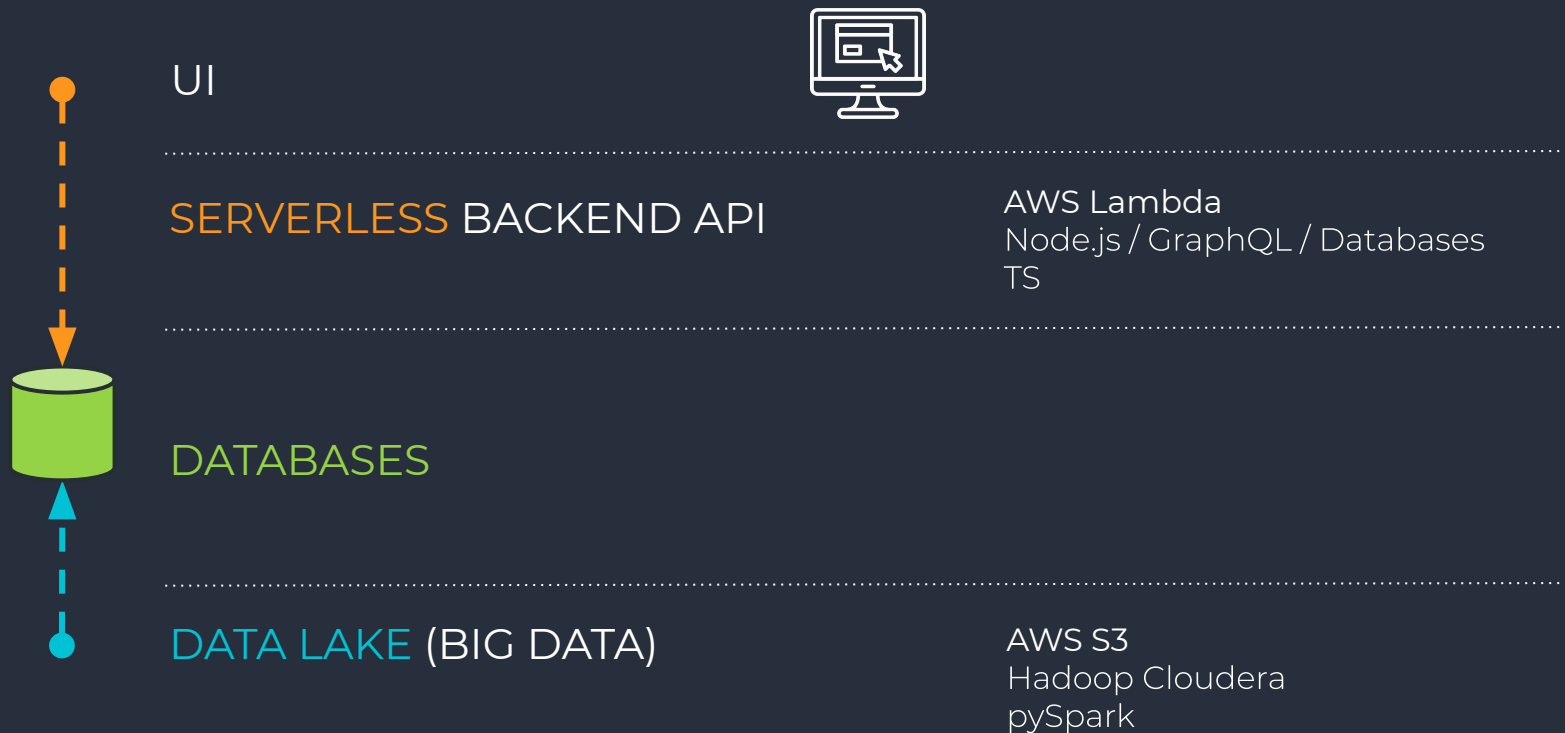
# PRODUCTION

## INTELLIGENCE SUPPLY CHAIN

- Hundreds GB RAW data and 60M new messages daily
- 8M UNIQUE ITEMS over the world

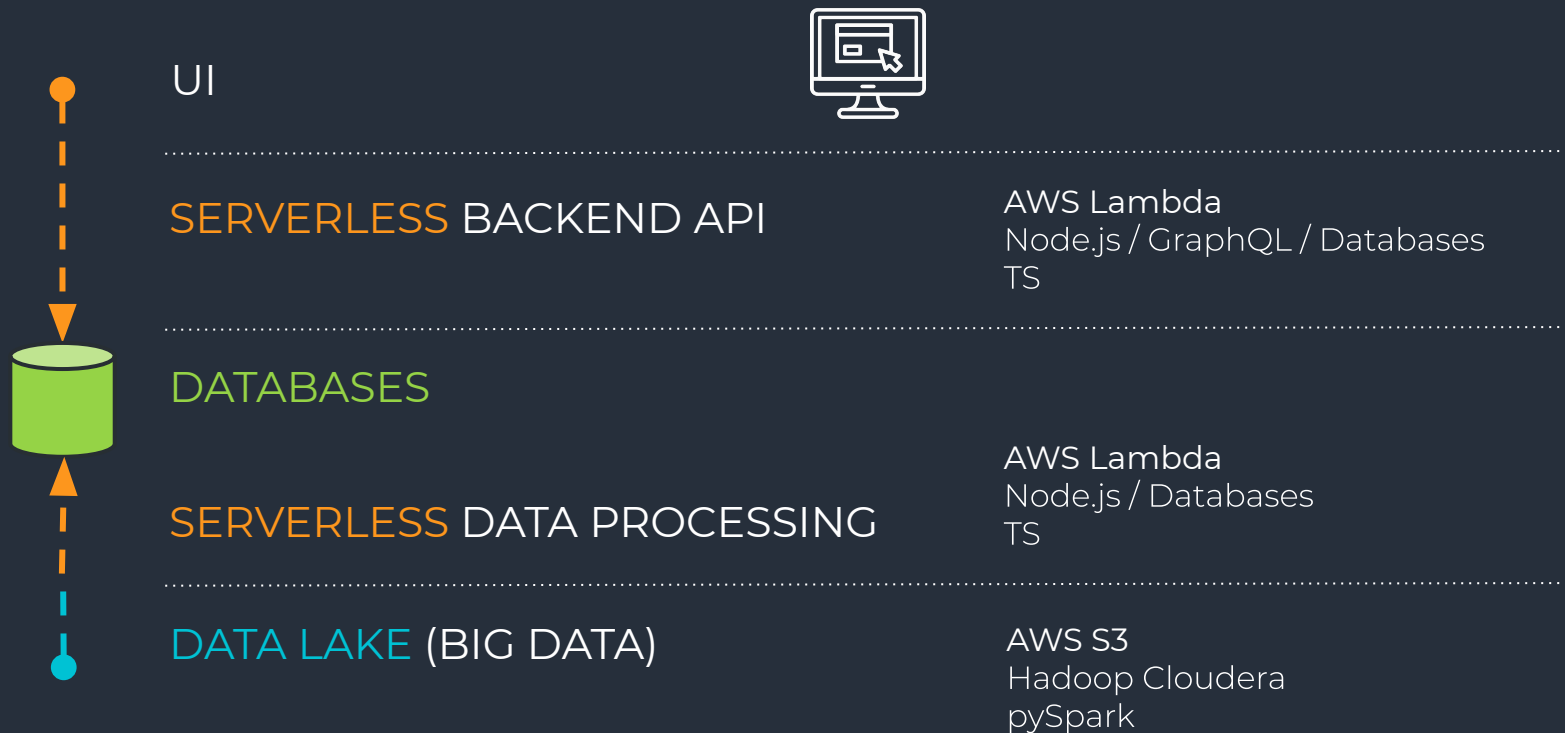


# TECHNOLOGIES





# TECHNOLOGIES



# SERVERLESS

“Focus on your core product instead of worrying about managing and operating servers or runtimes”



DATABASES  
DATA STORES



QUEUES,  
STREAMS,  
NOTIFICATIONS,  
EVENT BUS



FUNCTIONS  
ANALYTICS  
JOBS

# SERVERLESS FUNCTION

```
export const handler = async (event) => {  
  const data = event.Records[0].body;  
  
  // - TRANSFORM data  
  // - WRITE to DB or  
  // - PUT TO QUEUE/STREAM/TOPIC  
  
  return 'success';  
};
```




- NO INFRASTRUCTURE TO MANAGE
- TRIGGERED BY EVENTS
- HIGH SCALABLE
- STATELESS
- COST-EFFECTIVE
- CHOOSE YOUR CODE LANGUAGE

# NODE.JS POPULARITY

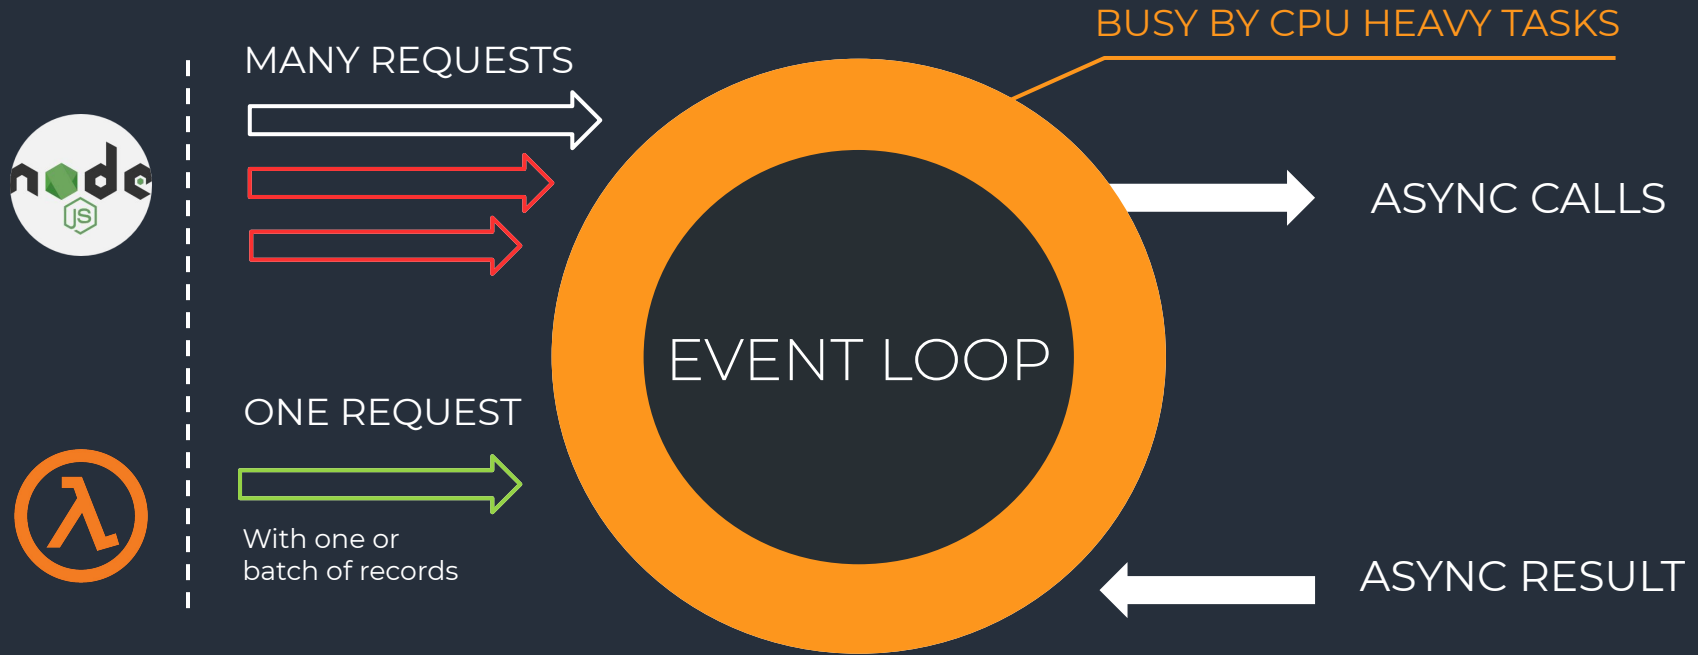
## 2014 AWS LAMBDA WITH NODE.JS

- PURE ASYNCHRONOUS
- MINIMALISTIC CORE
- FAST STARTUP WITH HIGH PERFORMANCE



2018 SERVERLESS.COM

# NODE.JS FUNCTION VS SERVER



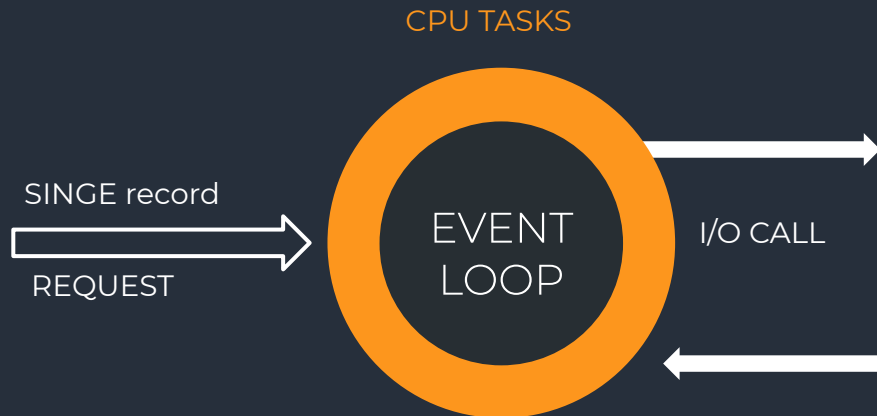
# I/O – PERFORMANCE BOTTLENECK



# SUMMARY

## CPU HEAVY TASKS


Get request with single record



---

## I/O INTENSIVE TASKS

Get request with batch of records (like a server)




# SERVERLESS COMPUTE SERVICE




THIS IS MY ARCHITECTURE – 177+ VIDEOS  
<https://aws.amazon.com/ru/this-is-my-architecture/>

Symbol in the presentation.

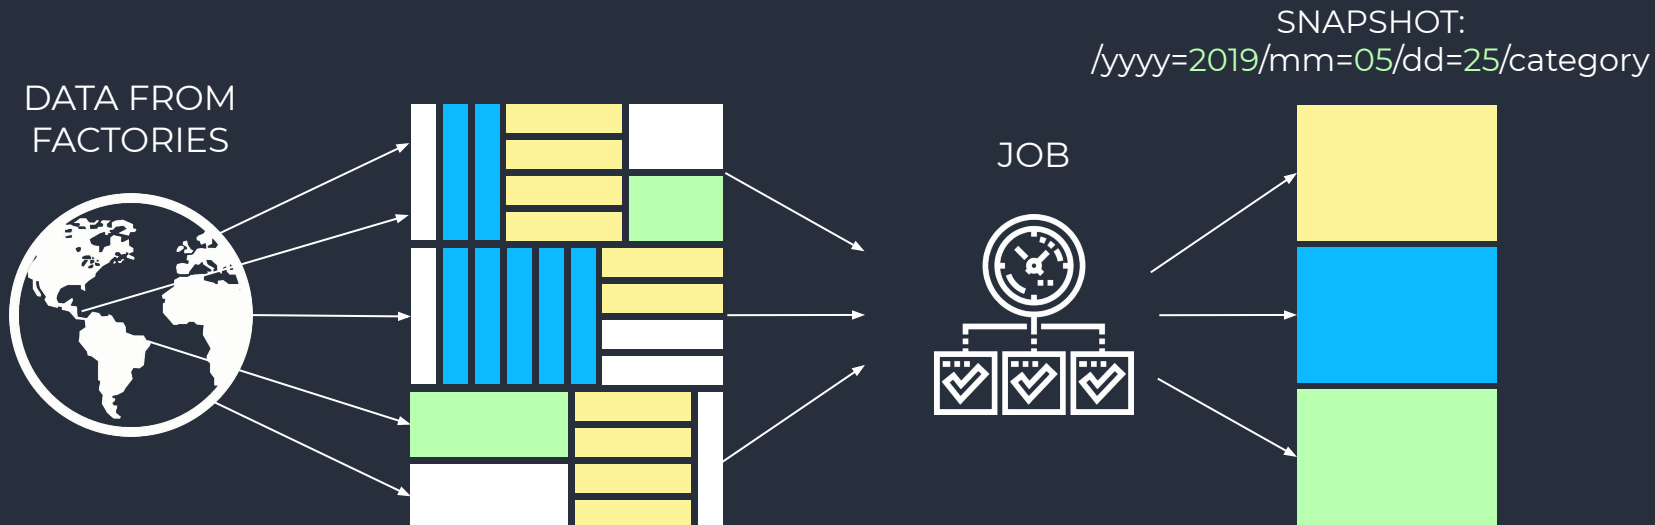




# EXAMPLE



# BATCH PROCESSING

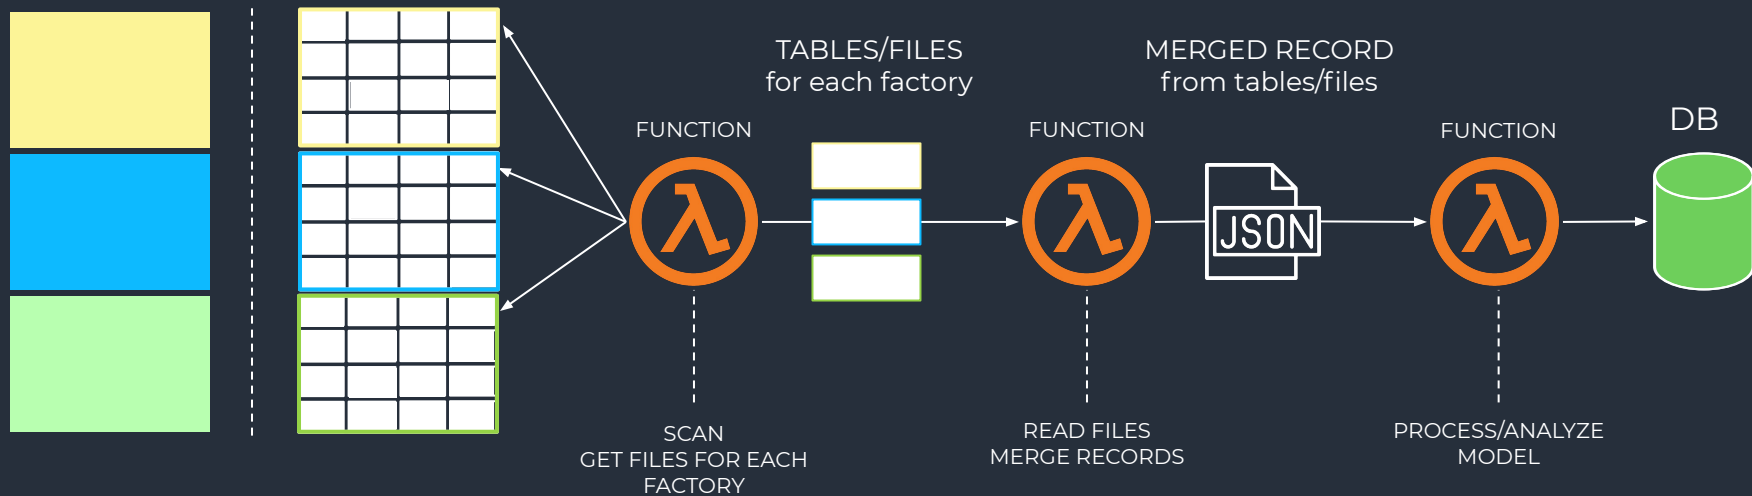


# BATCH DATA PROCESSING

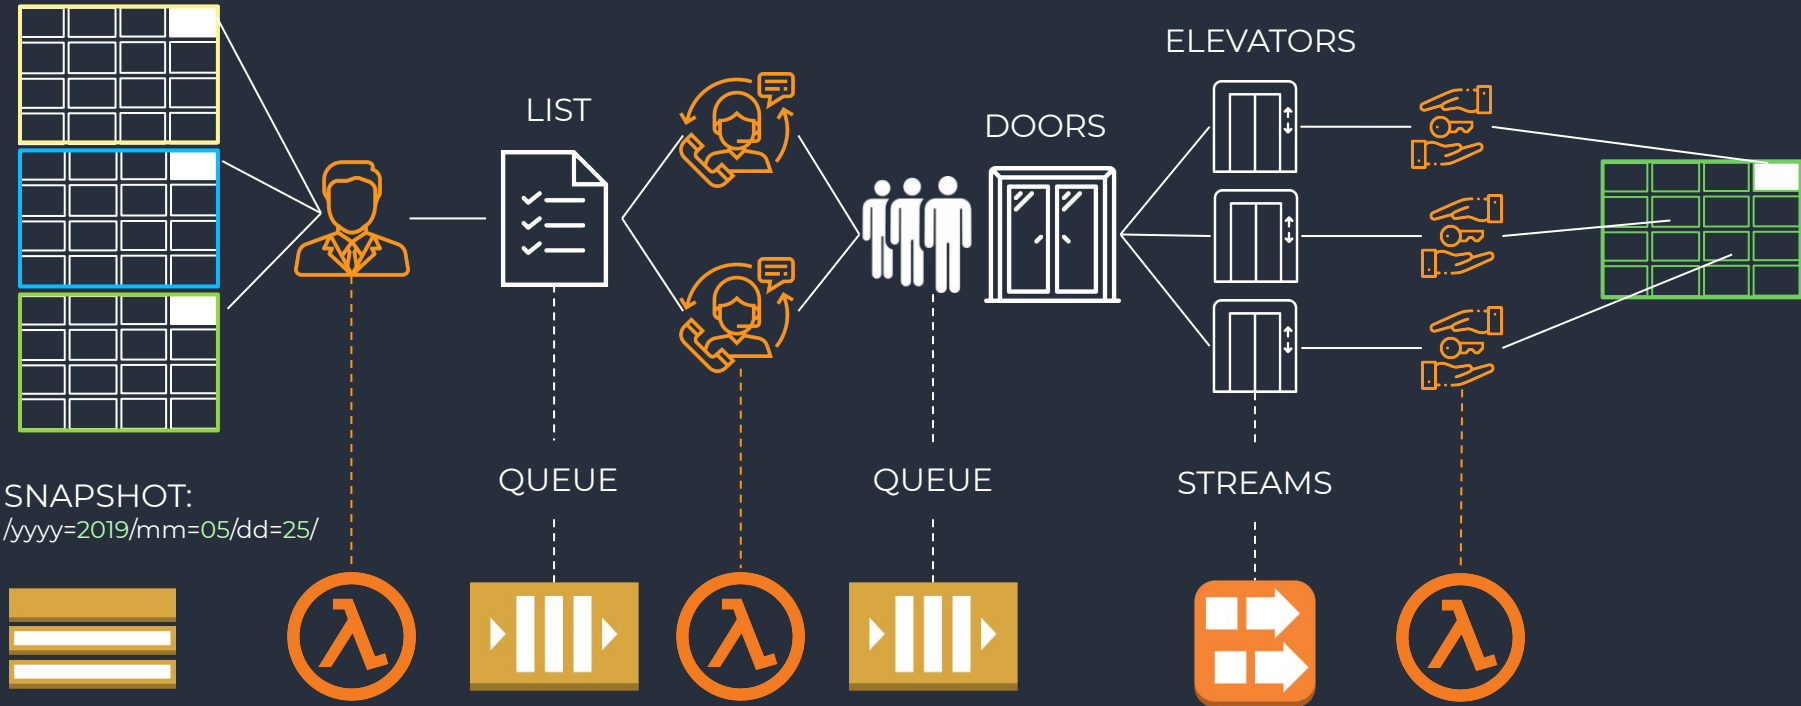
SNAPSHOT:

/yyy=2019/mm=05/dd=25/

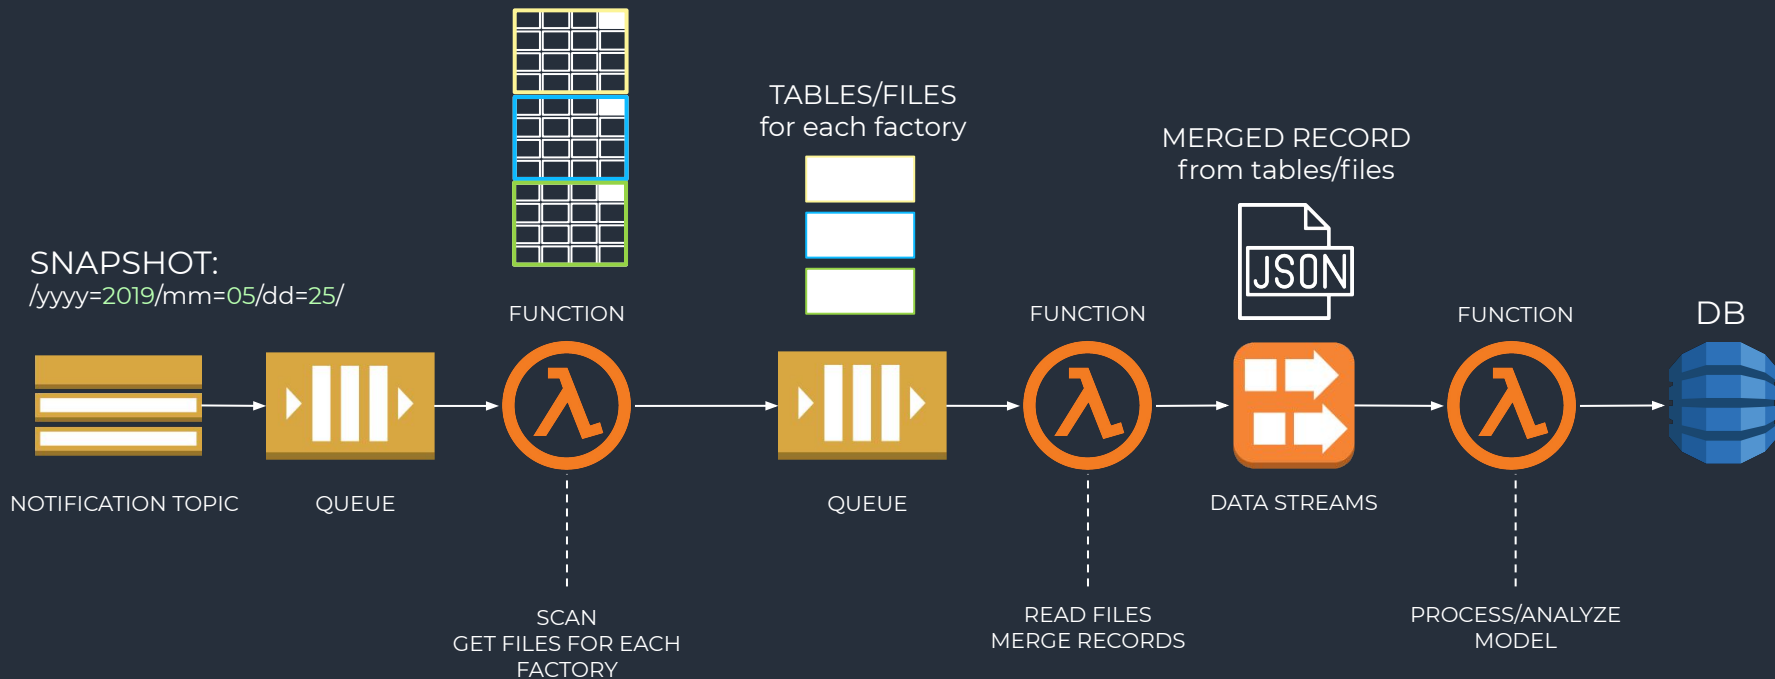
DONE!



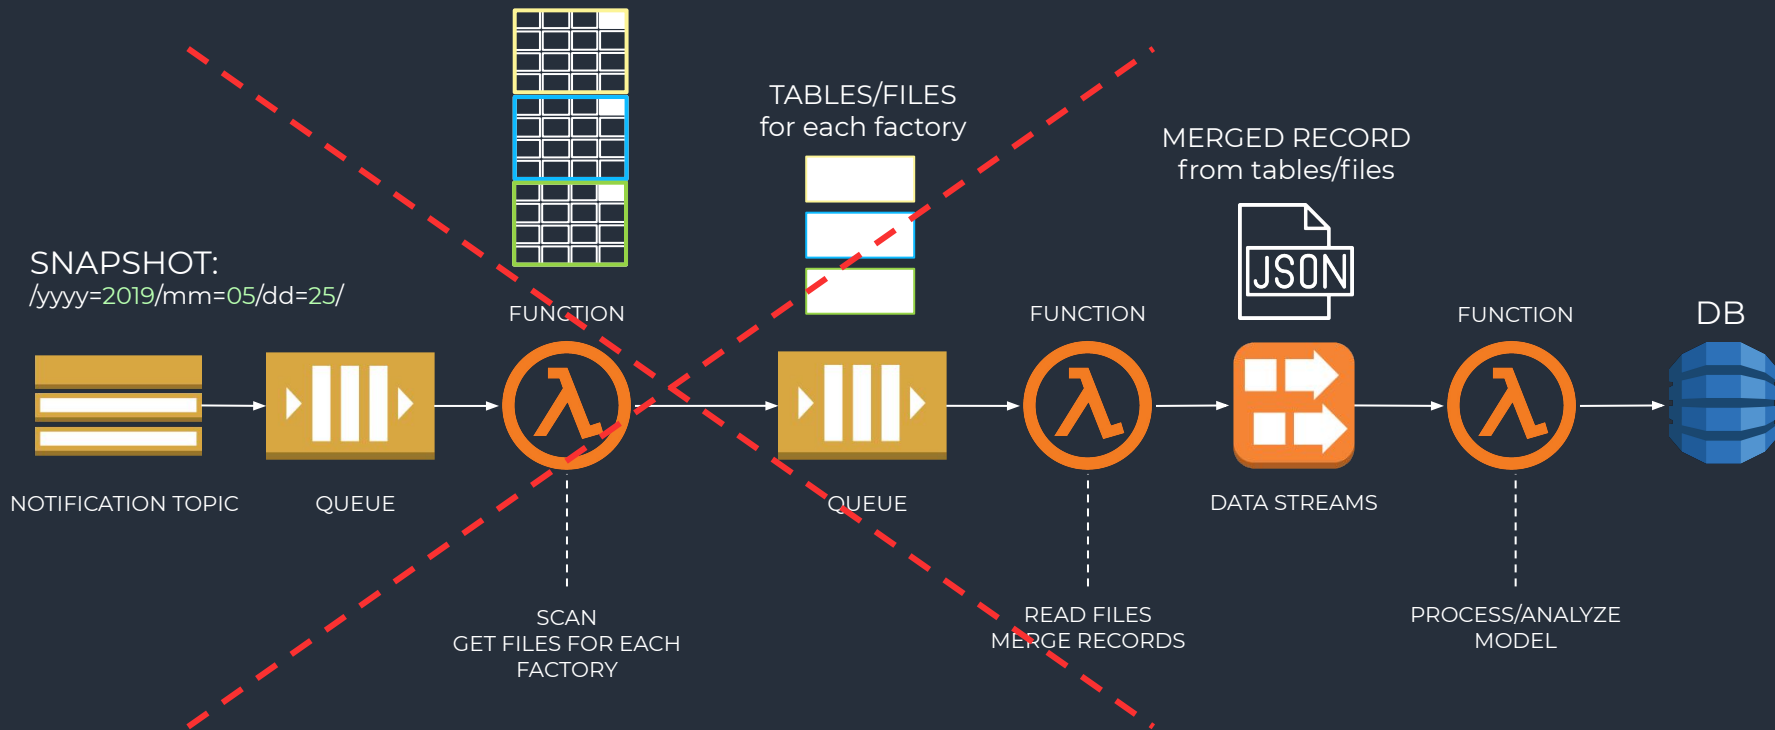
# RELOCATION OF PEOPLE TO A NEW BUILDING



# BATCH PROCESSING ARCHITECTURE



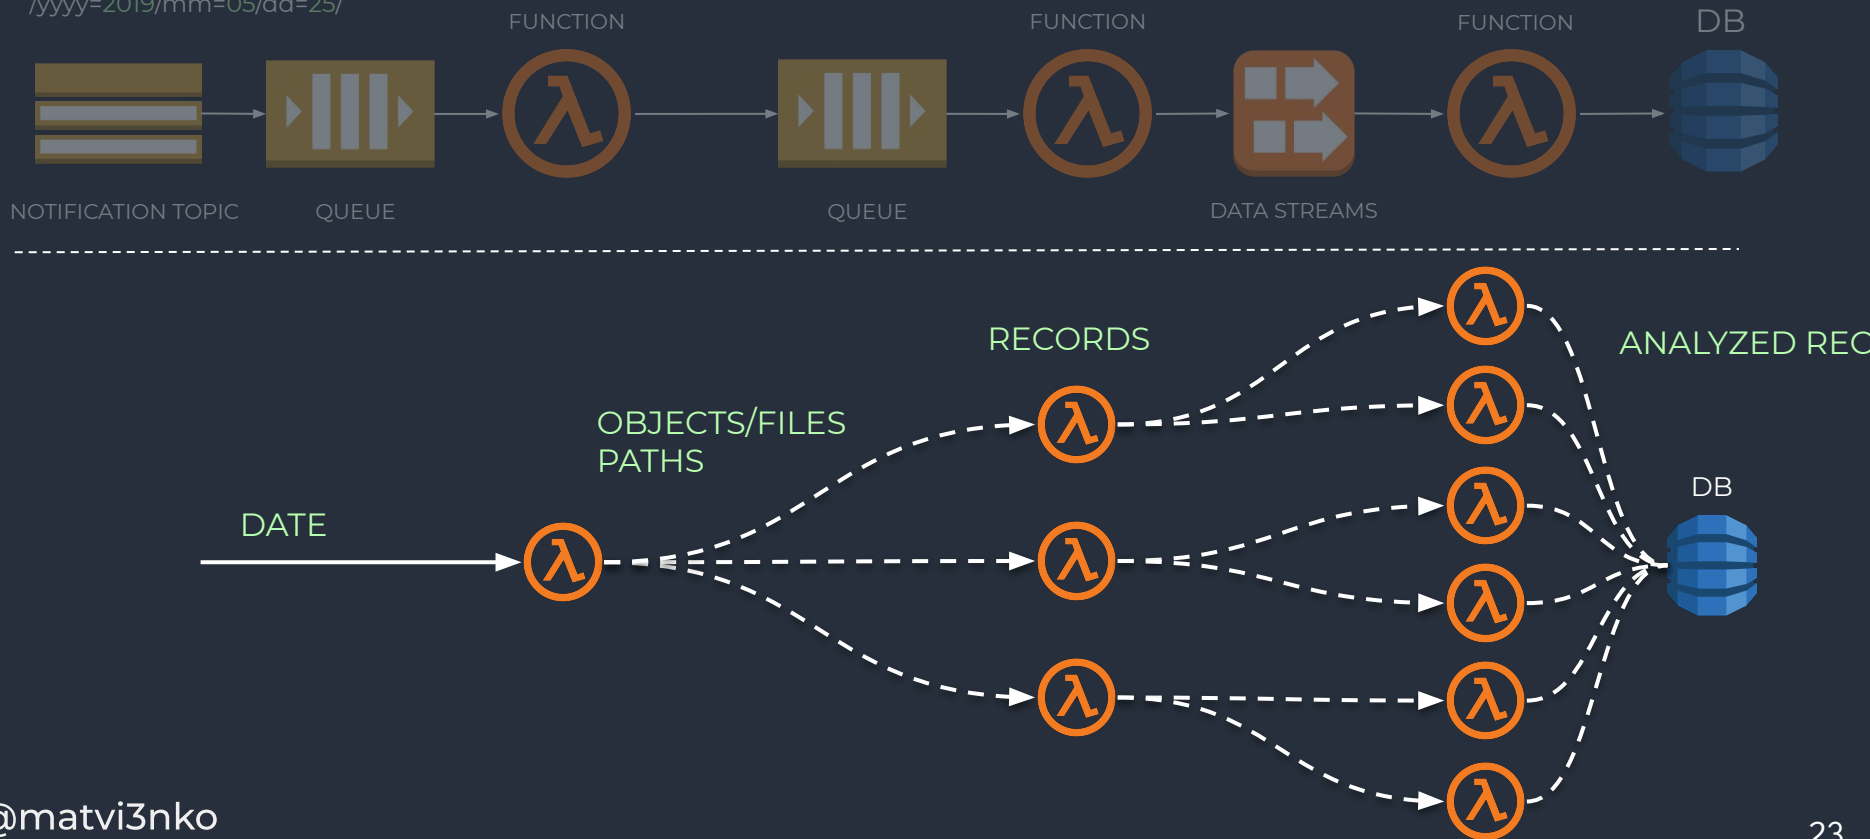
# REAL-TIME PROCESSING ARCHITECTURE



# SCALABILITY

SNAPSHOT:

/yyyy=2019/mm=05/dd=25/



# CONCLUSION

1. FROM BIG DATA TO A LARGE NUMBER OF MESSAGES
2. THE MORE MESSAGES THE FUNCTION ACCEPTS, THE MORE IT NEEDS TO BE PARALLELIZED
3. USE THE QUEUE FOR MESSAGES, AND DATA STREAMS TO TRANSFER MODELS / LARGE COLLECTION
4. INCREASE THE NUMBER OF STREAM SHARDS. 1 SHARD = 1 LAMBDA FUNCTION
5. PREPARE TO STREAMING / REAL-TIME PROCESSING



# PROGRESS



SOLUTION STRUCTURE AND FUNCTION

BASE ARCHITECTURE DESIGN

# SERVERLESS PROJECT STRUCTURE

/transform

– serverless.yml

– handler.ts

/analyze

– serverless.yml

– handler.ts

/node\_modules

serverless.yml

package.json

```
import AWS from 'as-sdk';
```

```
const s3Client = new AWS.S3({region});
```

```
export const handler = (event) => {
```

```
  const [message] = event.Records;
```

```
  return new Promise((resolve, reject) => {
```

```
    this.s3Client.selectObjectContent({ Key: message.path }, (err, data) => {
```

```
      if (err) {
```

```
        reject(err);
```

```
      }
```

```
      resolve(data);
```

```
    });
```

```
  };
```

# SERVERLESS PROJECT STRUCTURE

## /transform

- serverless.yml
- handler.ts

## /analyze

- serverless.yml
- handler.ts

## /node\_modules

serverless.yml  
package.json

## DISADVANTAGES

1. NODE\_MODULES contains dependencies of all functions  
Have to control and split them in SERVERLESS.YML
2. Lack of function isolation
3. Lack of independent install / build / test
4. Becomes monolith project

# MONOREPO SERVERLESS PROJECTS STRUCTURE

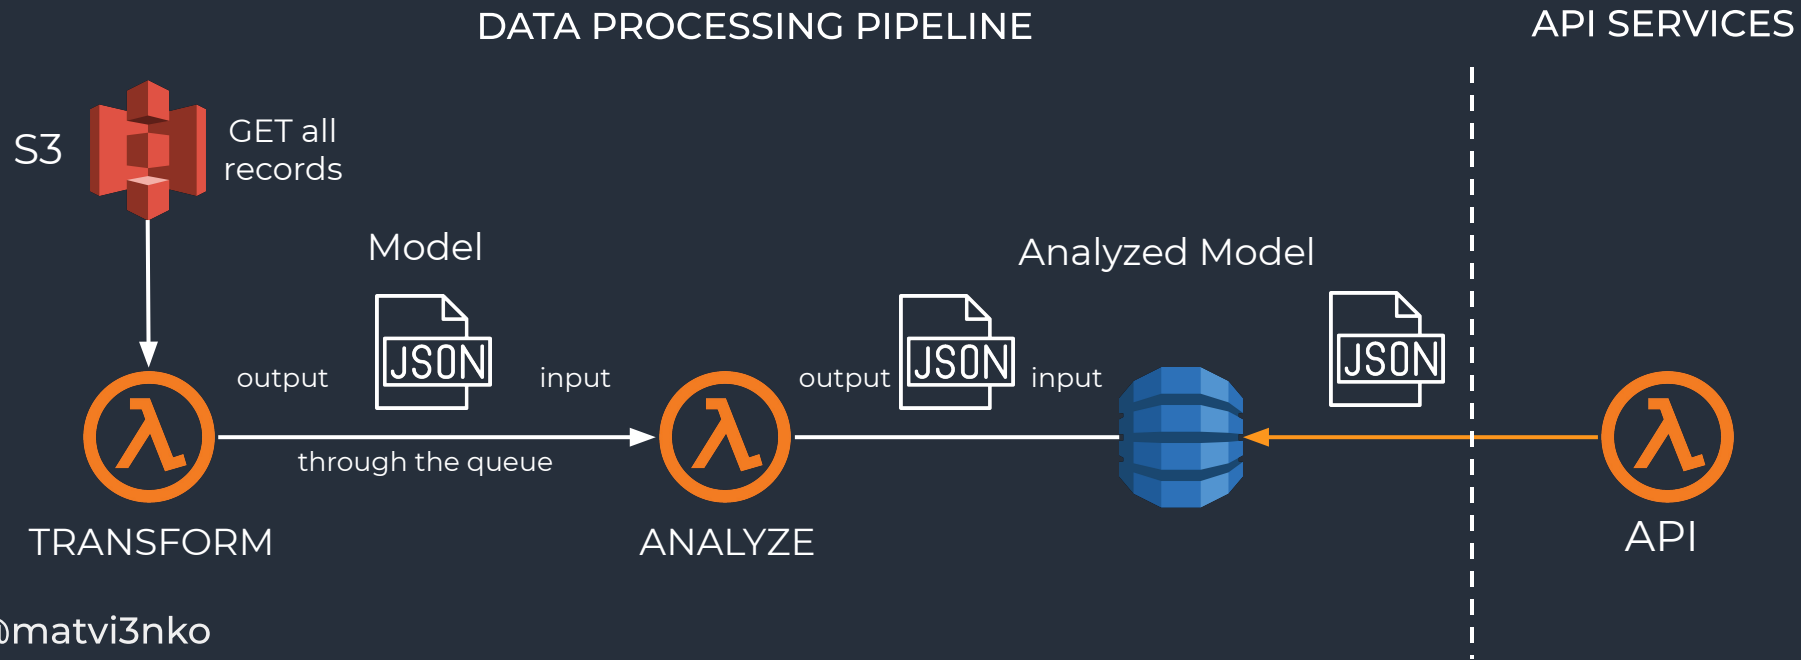
```
/lib
  /node_modules
  /errors
  /factories
  /models
  /providers
- package.json
/transform(er)
  /node_modules
- functionA1.ts
- package.json
- serverless.yml
/analyze(r)
  /node_modules
- functionA2.ts
- package.json
- serverless.yml
```

## ADVANTAGES

1. LIB contains all common infrastructure, domain logic and cloud provider's SDK
2. Functions became isolated projects with flexible splitting and contains only business logic
3. LIB and PROJECTS versioning
4. NPM resolves NODE\_MODULES dependencies automatically
5. Independent install / build / test / deploy / troubleshooting

# MODELS REUSABILITY

```
import { Model } from '@holyls/models'
```



# 3 MAIN PRINCIPLES OF ANY FUNCTION

1. INITIALIZE ONLY ONCE
2. USE STREAMING PROCESSING
3. HANDLE ERRORS AT A HIGH LEVEL


# RECOMMENDED

```
const s3 = createCsvS3Provider(config.region);  
const streams = createKinesisProvider(config.region);  
const service = new SomeService();
```

```
const handler = async (event) => {  
  const source = s3.getObjects(event.Records[0].body.path)  
  .pipe(flatMap(([a, b, c]) => from(service.transform(a, b, c)))  
  .pipe(flatMap(model => service.compress(model)))  
  .pipe(bufferCount(config.batchSize))  
  .pipe(flatMap(batch => streams.put(batch)));
```

```
return new Promise((resolve, reject) => {  
  source.subscribe(() => { /* handle */,
```

```
    err => {  
      err instanceof InfrastructureError && reject(err);  
      err instanceof DomainError && reject(err);  
    }, resolve);});};
```




# NOT RECOMMENDED

```
const handler = async (event) => {  
  const s3 = createCsvS3Provider(config.region);  
  const streams = createKinesisProvider(config.region);  
  const service = new SomeService();
```

```
  const objects = await s3.getObjects(event.Records[0].body.path);  
  const models = await service.transform(objects);
```

```
  return Promise.all(models.map(model => {  
    const compressedModel = await service.compress(model);  
    return streams.put(compressedModel);  
  }));  
};
```



# 4 COMPONENTS OF THE FUNCTION

1. INCOMING DATA
2. TYPE OF TRANSFORMATION
3. DESTINATION
4. MAIN ERROR HANDLER



# PROGRESS




ERROR HANDLING

SOLUTION STRUCTURE AND FUNCTION

BASE ARCHITECTURE DESIGN

# TRACK LOG MESSAGES THAT ARE HIDDEN ERRORS

1. Request XX-YY: "Process exited before completing request"
2. Function completed on its timeout (up to limit)



# HOW TO HANDLE

## 1. FUNCTION HANGS

Don't do that: `context.callbackWaitsForEmptyEventLoop = false;`

or close Sequelize connection to fix that

Use callback cb(), rewrite to async/await (Promises)

## 2. FUNCTION DOES NOT PERFORM PART OF THE LOGIC

You added async or return value.

Find missed await / return Promise



# WAIT FOR RESPONSES FROM THE SERVICES

// CODE

SEND(MESSAGE);

// CODE

// CODE

RETURN SEND(MESSAGE);

// CODE

// CODE


AWAIT SEND(MESSAGE);

// CODE

SEND



SAVE 1\$



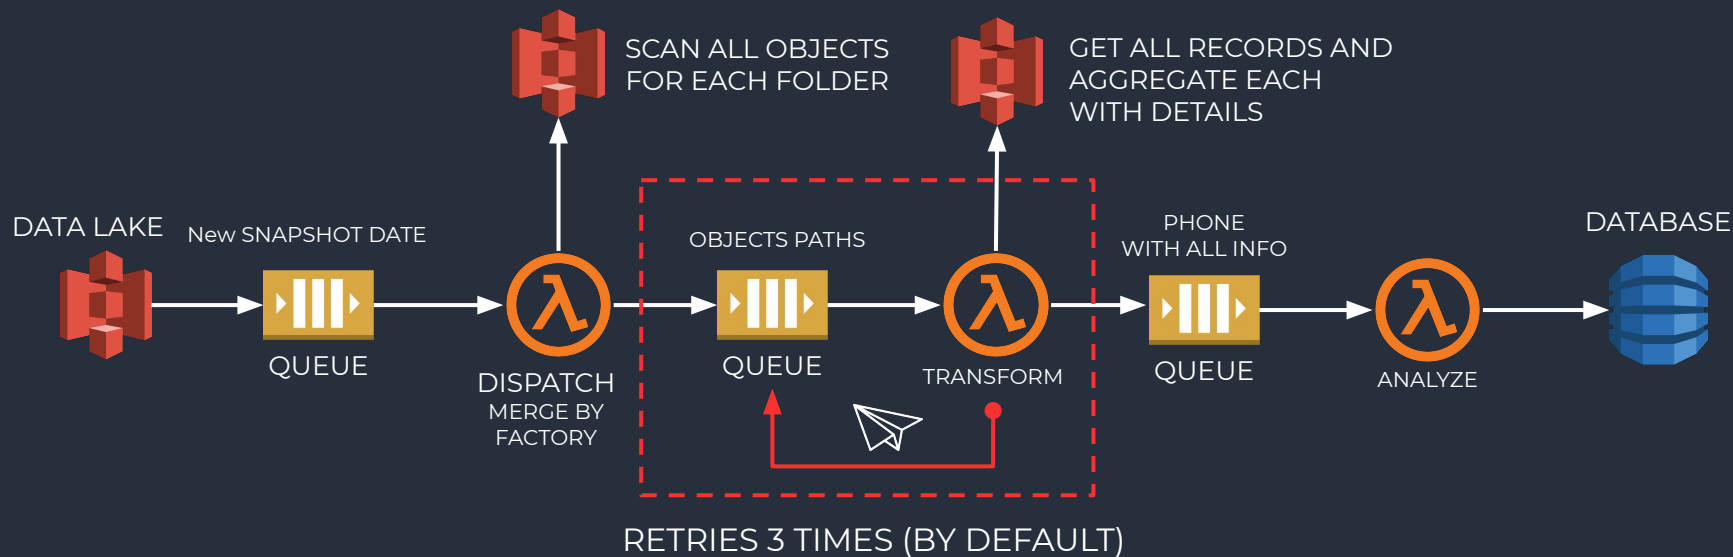
LOSE MILLIONS \$



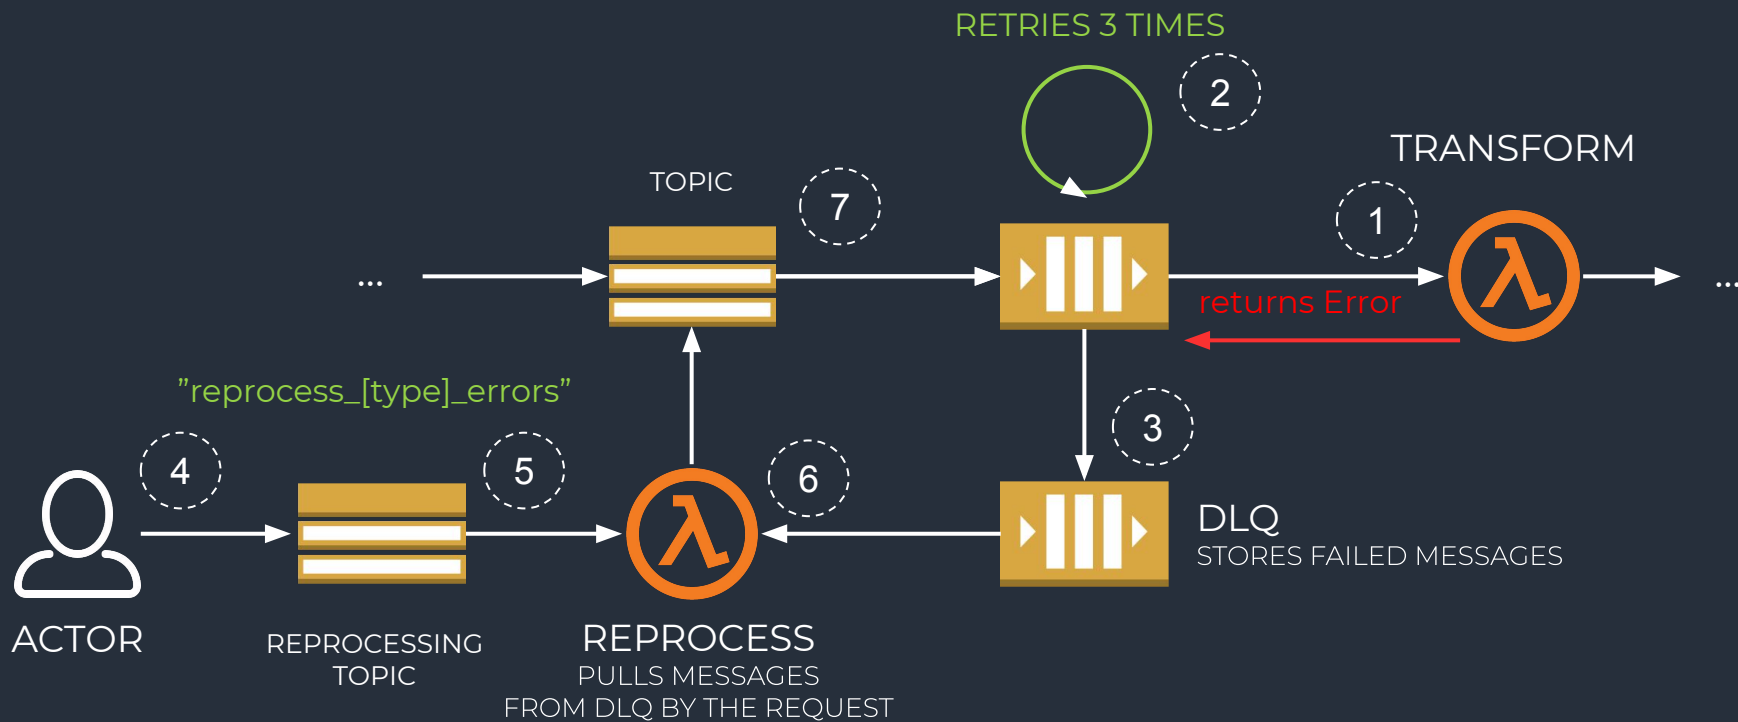
ERROR!



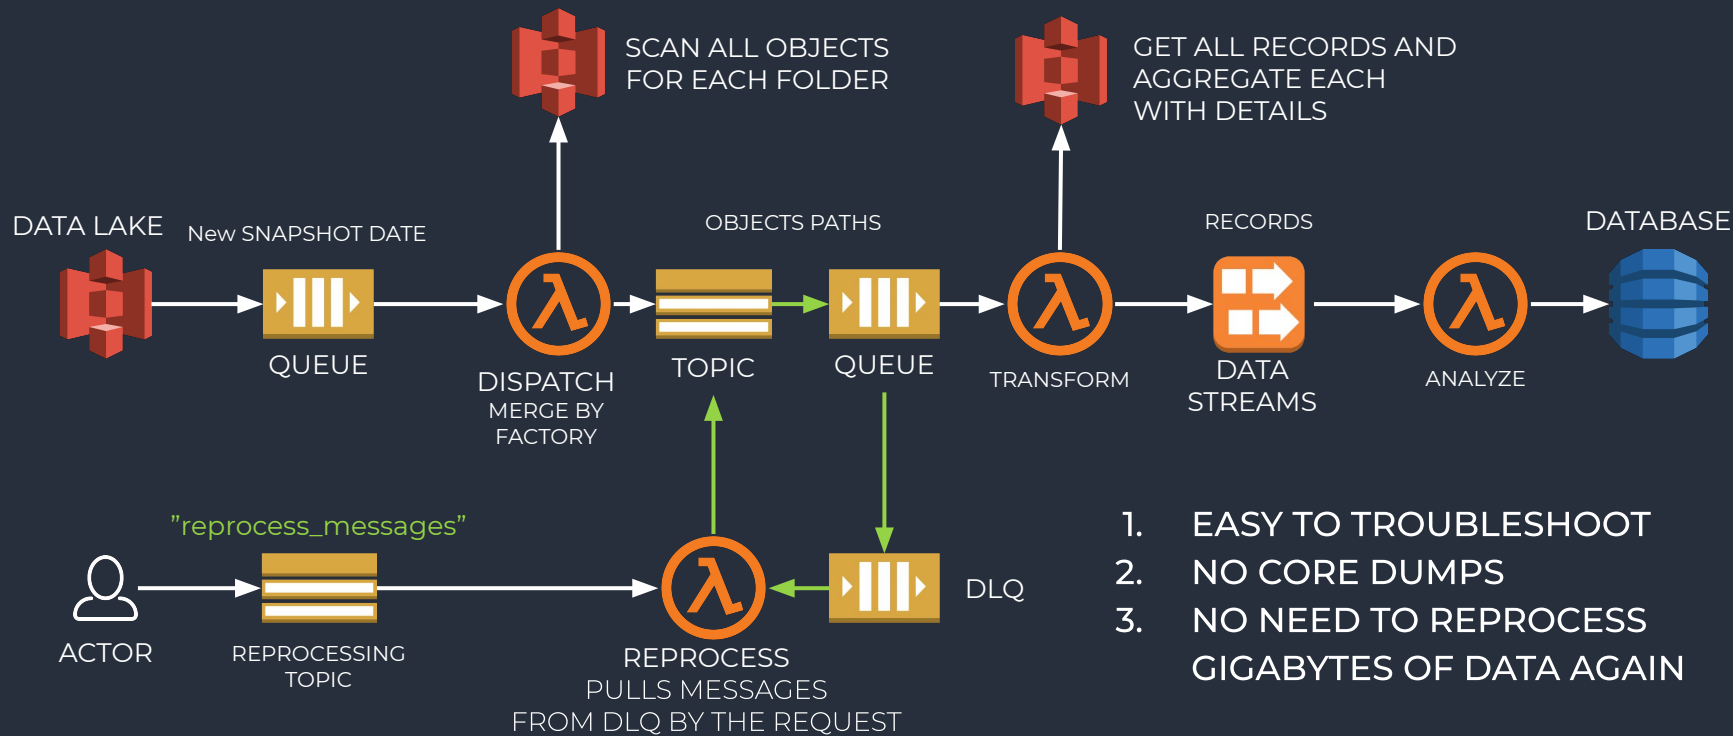
# RETRY STRATEGY



# DEAD LETTER QUEUE

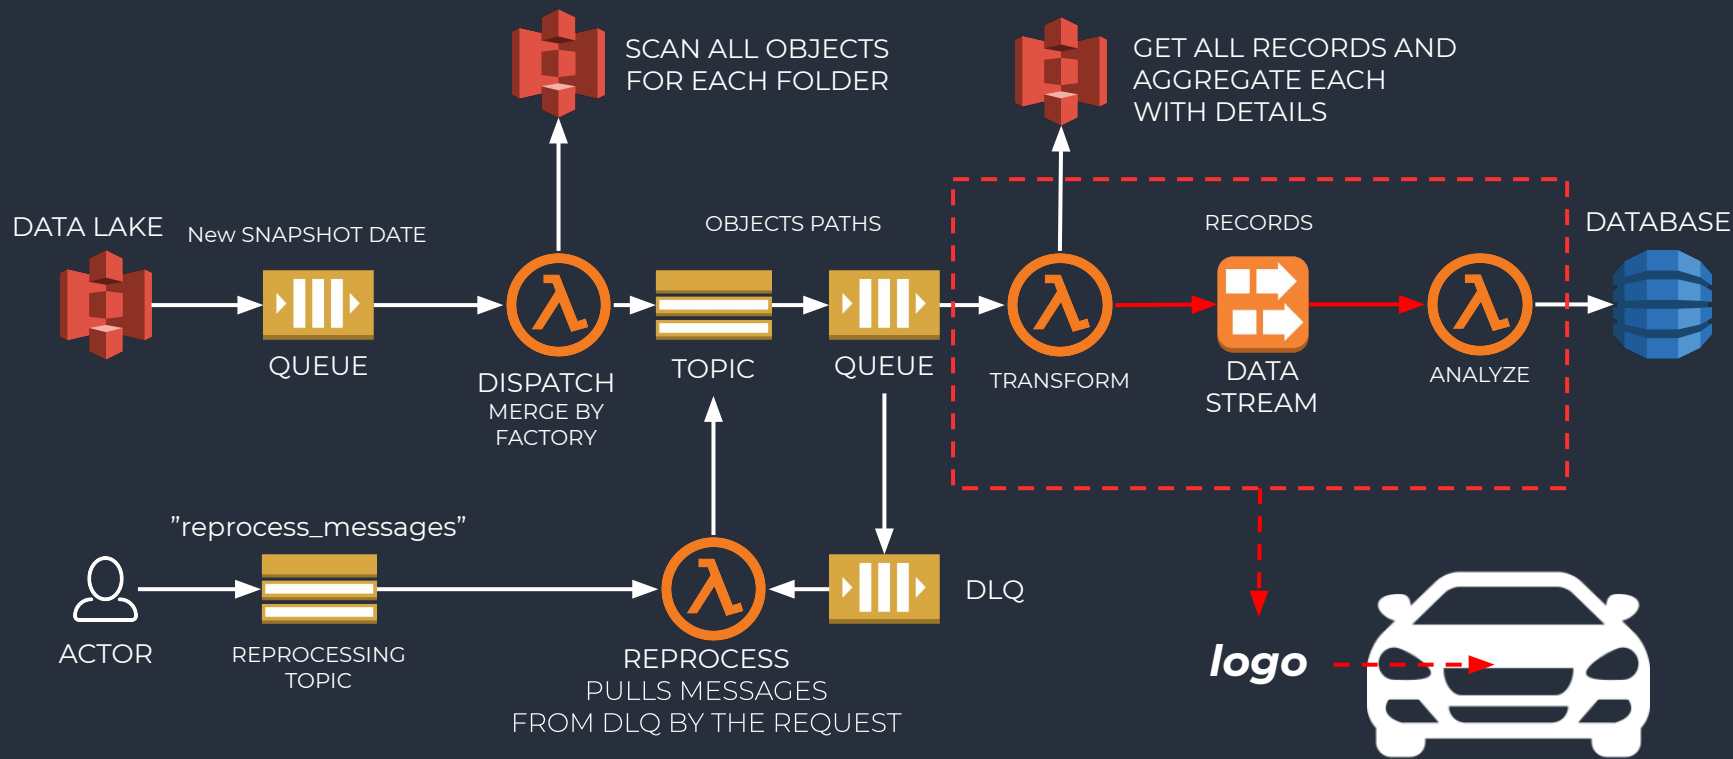


# DLQ FOR THE QUEUE



1. EASY TO TROUBLESHOOT
2. NO CORE DUMPS
3. NO NEED TO REPROCESS GIGABYTES OF DATA AGAIN

# KINESIS ERROR HANDLING





# ERROR HANDLING STRATEGY

## 1. PUT TO STREAM (TRANSFORM)

- a. Handle partially successful response

```
FailedRecordCount : Number
```



## 2. READ FROM (ANALYZE)

- a. Reduce retry times


```
customBackoff: (retryCount) => {...}
```

- b. Use Dead Letter Queue

- c. Handle duplicate Records



UP TO 7 days each 100 ms



# EXACTLY ONCE PROCESSING

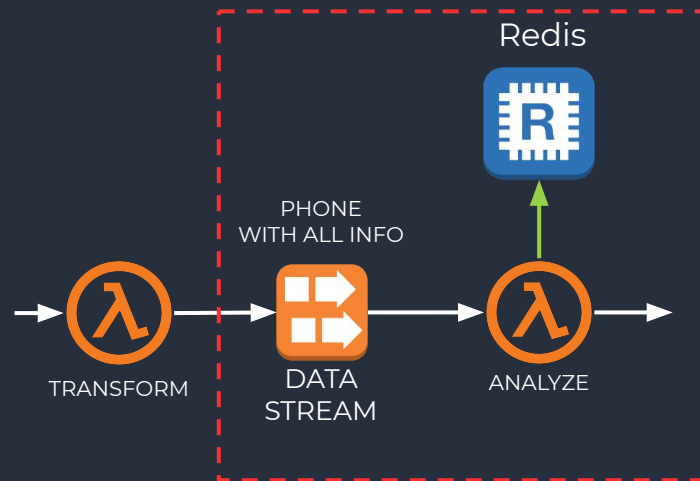
1. REDUCE THE RISK OF FAILURE

```
const response = await putToStream(record);  
// do something with response -> RISK
```

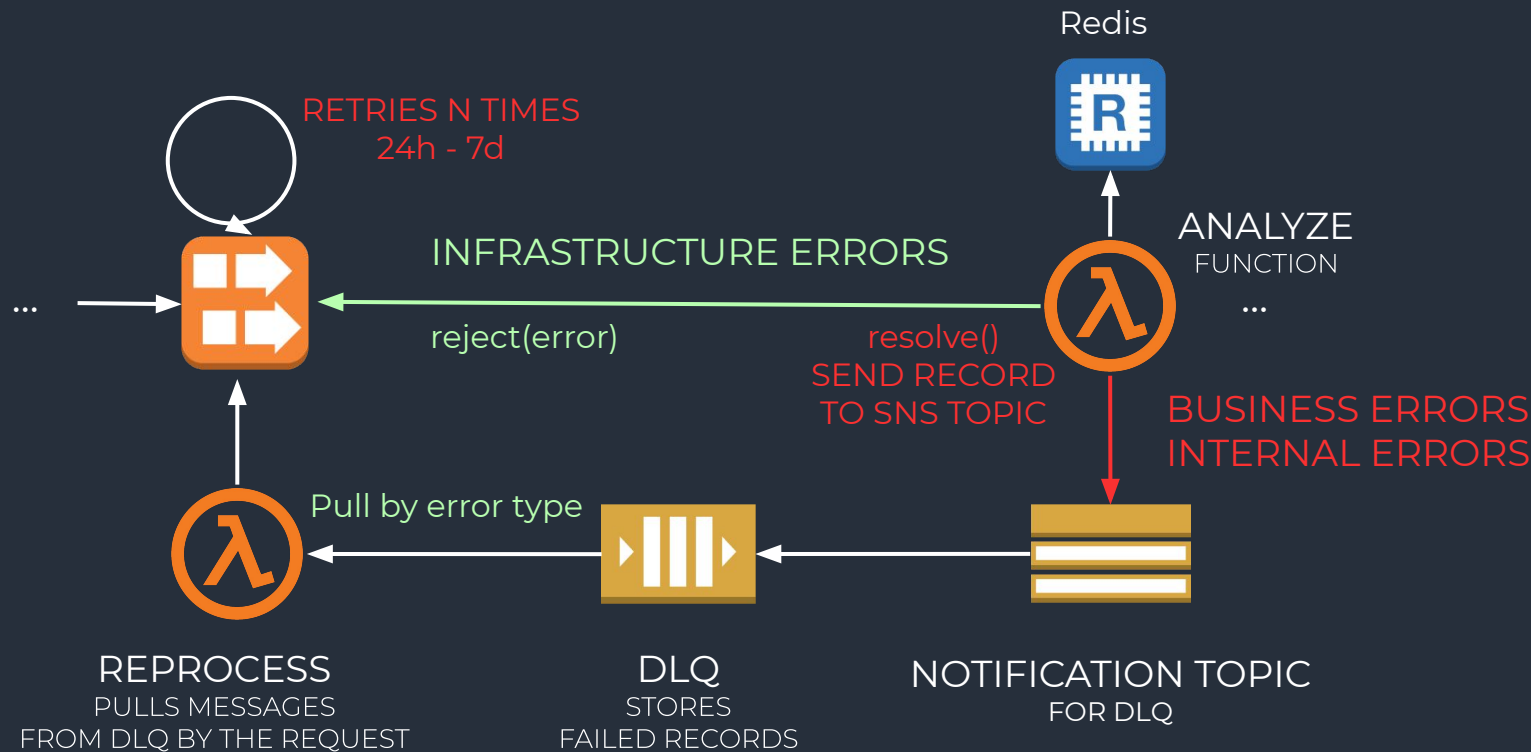
2. USE REDIS CACHE TO STORE KEYS OF RECORDS

key: [date]-[shard-id]-[sequence-number]

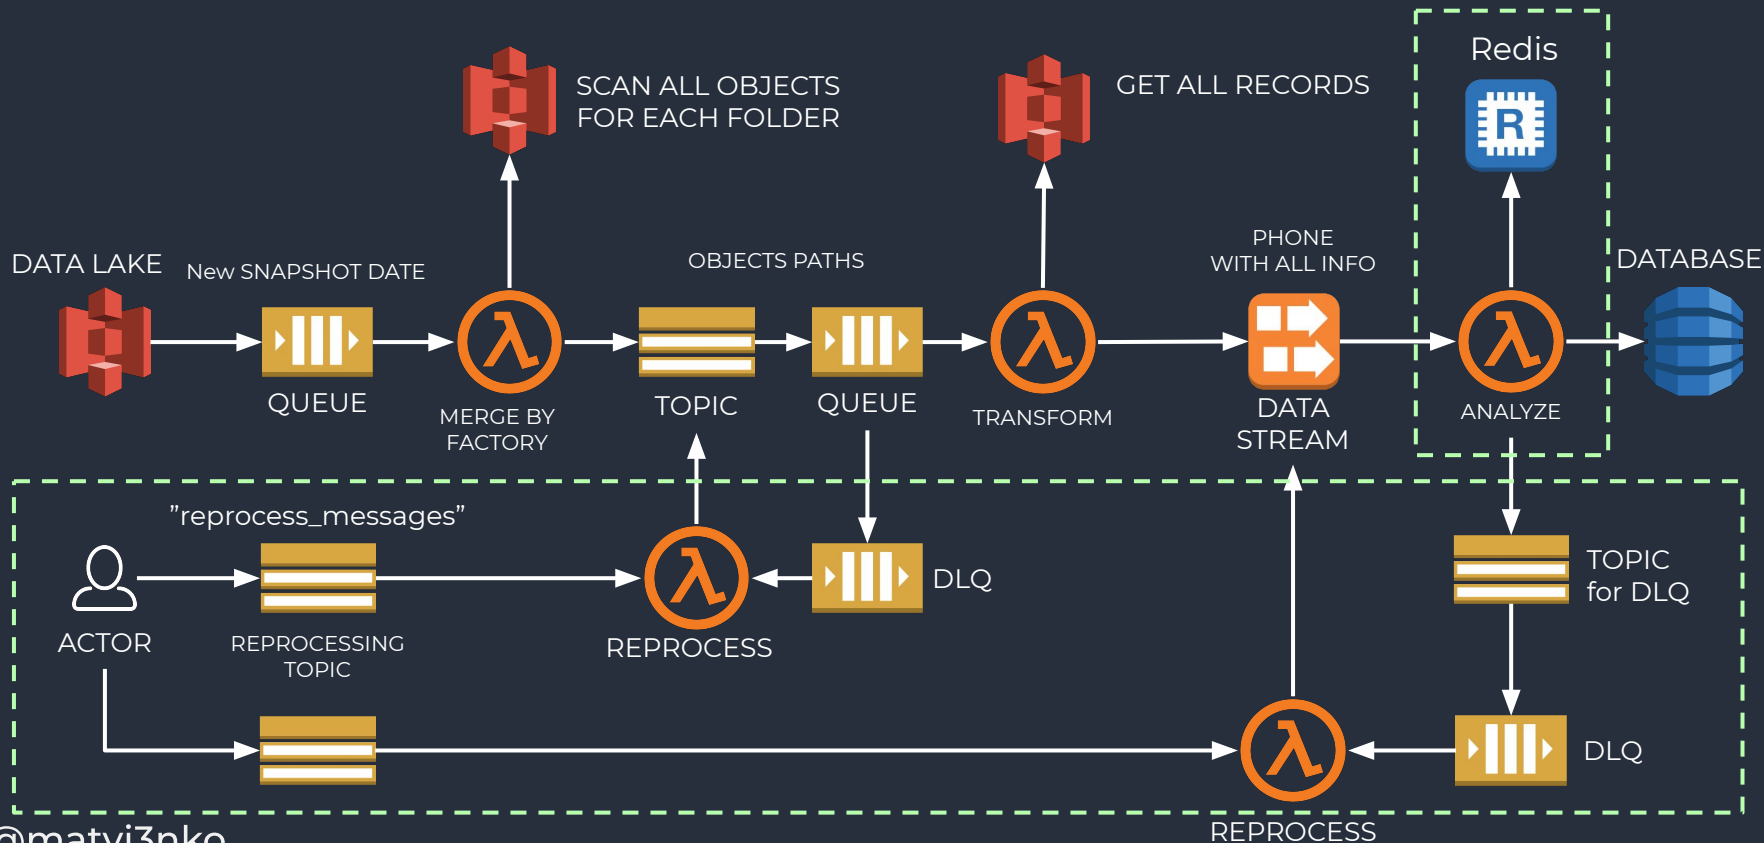
3. FILTER DUPLICATES



# KINESIS ERROR HANDLING



# ARCHITECTURE WITH REPROCESSING BLOCK



# CONCLUSION

1. CREATE CUSTOM TYPES OF ERRORS
2. REPROCESS ONLY FAILED MESSAGE AND NOT GIGA/TERABYTES OF DATA
3. DON'T LOSE MESSAGES
4. USE DLQ WITH ERROR FILTERING
5. PROCESS EXACTLY ONCE

# PROGRESS



READ AND PROCESS DATA

ERROR HANDLING

SOLUTION STRUCTURE AND FUNCTION

BASE ARCHITECTURE DESIGN

# AWS S3: SIMPLE STORAGE SERVICE

## 1. S3 SELECT REQUEST

```
const params = {  
  Bucket: 'bucket_name',  
  Key: 'key_name',  
  ExpressionType: 'SQL',  
  Expression: `  
    SELECT s.name, s.year  
    FROM S3Object s  
    WHERE s.name = 'HolyJS'  
  `,  
},  
OutputSerialization: {  
  JSON: { RecordDelimiter: '\n' },  
};
```



## 2. QUERY RESULT

id	name	year	details

JSON | CSV | Parquet

UP TO 4x FASTER

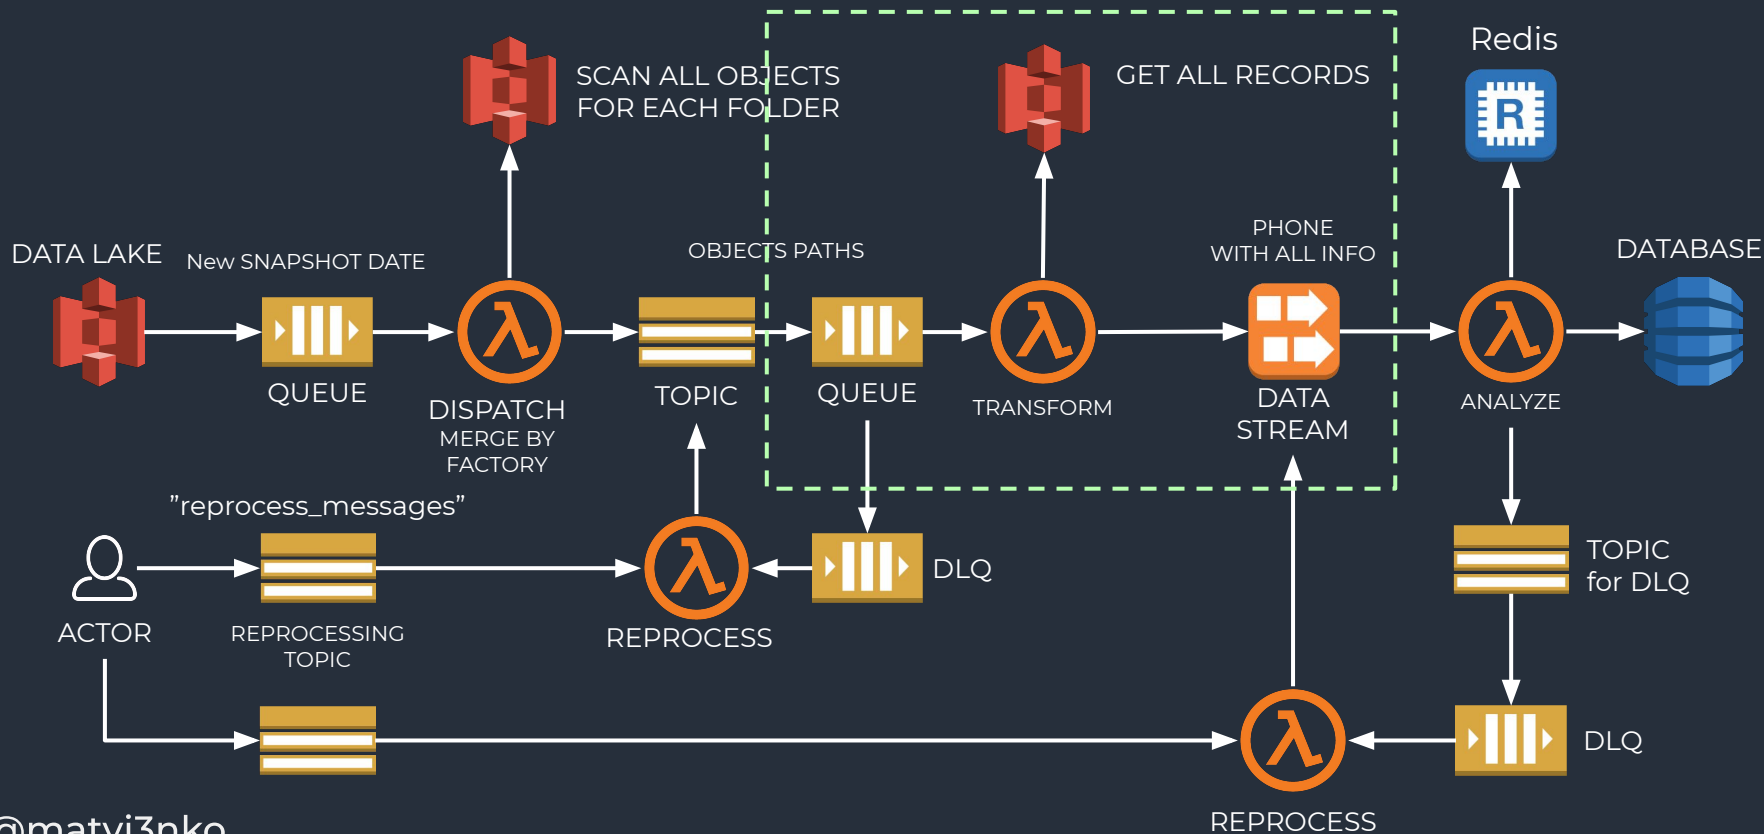
# S3: CSV VS PARQUET

Dataset	Size on Amazon S3	Query Run time	Data Scanned	Cost
Data stored as CSV files	1 TB	236 seconds	1.15 TB	\$5.75
Data stored in Apache Parquet format*	130 GB	6.78 seconds	2.51 GB	\$0.01
Savings / Speedup	87% less with Parquet	34x faster	99% less data scanned	99.7% savings

<https://dzone.com/articles/how-to-be-a-hero-with-powerful-parquet-google-and>



# EXTRACTING BIG OBJECTS/FILES



# DATA FORMATS

S3 OBJECT PATH:

s3/buckets/bucket-name/**entity**/yyyy=2019/mm=05/dd=25/**partition-by-category**/key.parquet



RAW OR COMPRESSED?



**80MB**  
PARQUET  
BINARY



**1.4GB**  
CSV  
TEXT. ONLY VALUES

**~X20 bigger size**



**6.5GB**  
JSON  
TEXT. FIELDS + VALUES

**~X80 bigger size**  
**75% - fields names**

# JSON

CHUNK 1 – 64KB

```
{
  sourceId: '1',
  description: 'device1',
  createdAt: "2019-02-01 00:00:00",
  someValue: "17",
  someSpecificTotalAmount: "13.50",
}, {
  source
```

CHUNK 2 – 64KB

```
    eld: '2',
  description: 'device2',
  createdAt: "2019-02-01 00:00:00",
  someValue: "18",
  someSpecificTotalAmount: "14.64",
},
```

# CSV

CHUNK 1 – 64KB

```
`"1","device1","2019-02-01 00:00:00","17","13.50"\n
"2","device2","2019-02-01 00:00:00","18","14.64"\n
"3","device3","2019-02-01 00:00:00","19","15.11"\n`
"4","device4","2019-02-" ...
```

CHUNK 2 – 64KB

```
...`"01 00:00:00","20","16.43"\n
"5","device5","2019-02-01 00:00:00","19","15.09"\n
"6","device5","2019-02-01 00:00:00","19","15.09"\n
"7","device5","2019-02-01 00:00:00","19","15.09"\n
```

# PARQUET

```
010101100111000110111011011
001110001101110110110011100
011011101101100111000110111
01
```

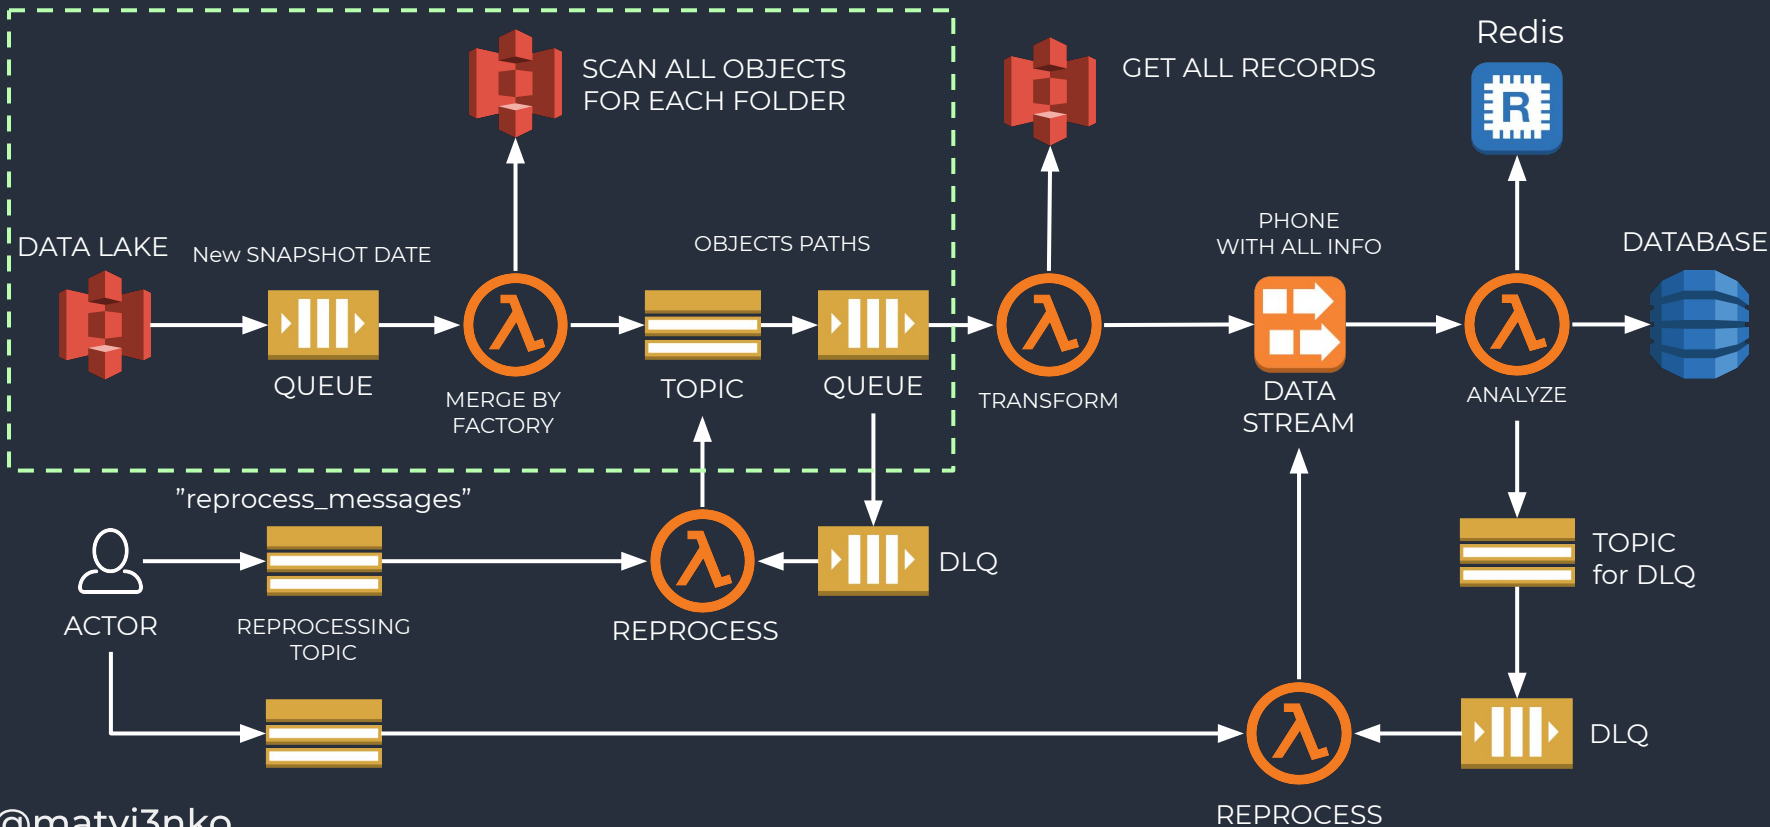
```
<Buffer 61 61 3a 20 34 2c 20 62
3a 20 35 2c 20 63 3a 20 34 2c 20
20 31 30 30 30>
```

+ DECODING IN STREAM

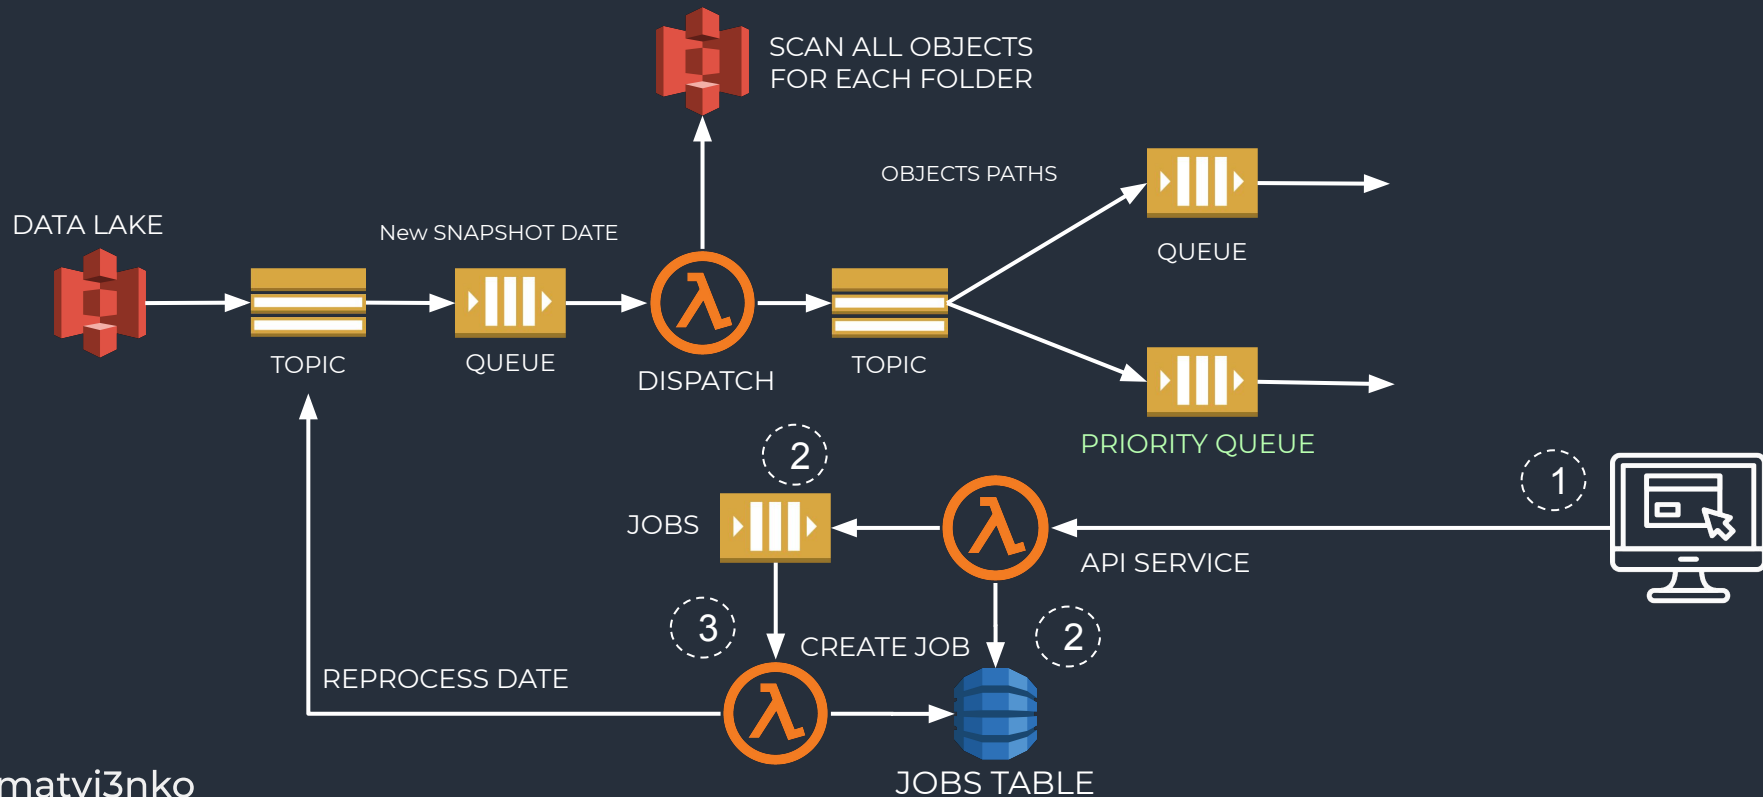
# CONCLUSION

1. DON'T USE JSON FOR READING/PROCESSING
  - a. Big size
  - b. Slow parsing
2. USE PARQUET (or CSV in some cases)
  - a. Parquet decoding Increases GC time (5 - 15% of time)
3. BUT DON'T BUFFER RESPONSE, WORK WITH STREAMS
4. GZIP RECORDS BEFORE PUT TO KINESIS STREAM
5. INCREASE MEMORY (CPU AUTO-LY) TO WORK WITH BIG OBJECTS/FILES

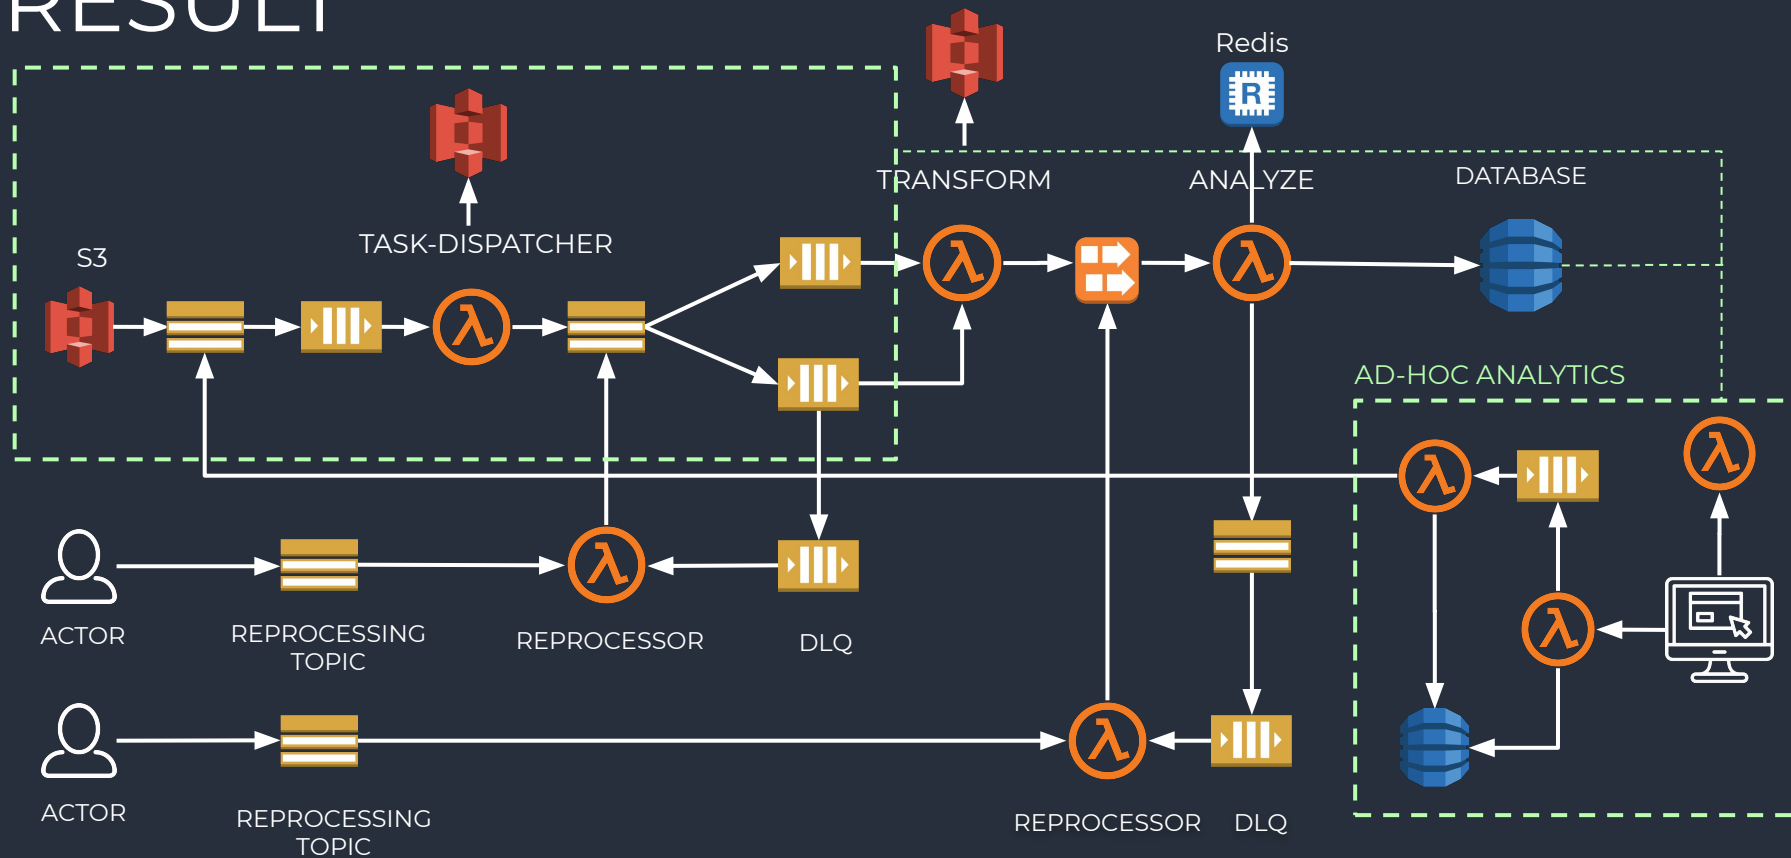
# AHEAD OF THE QUEUE



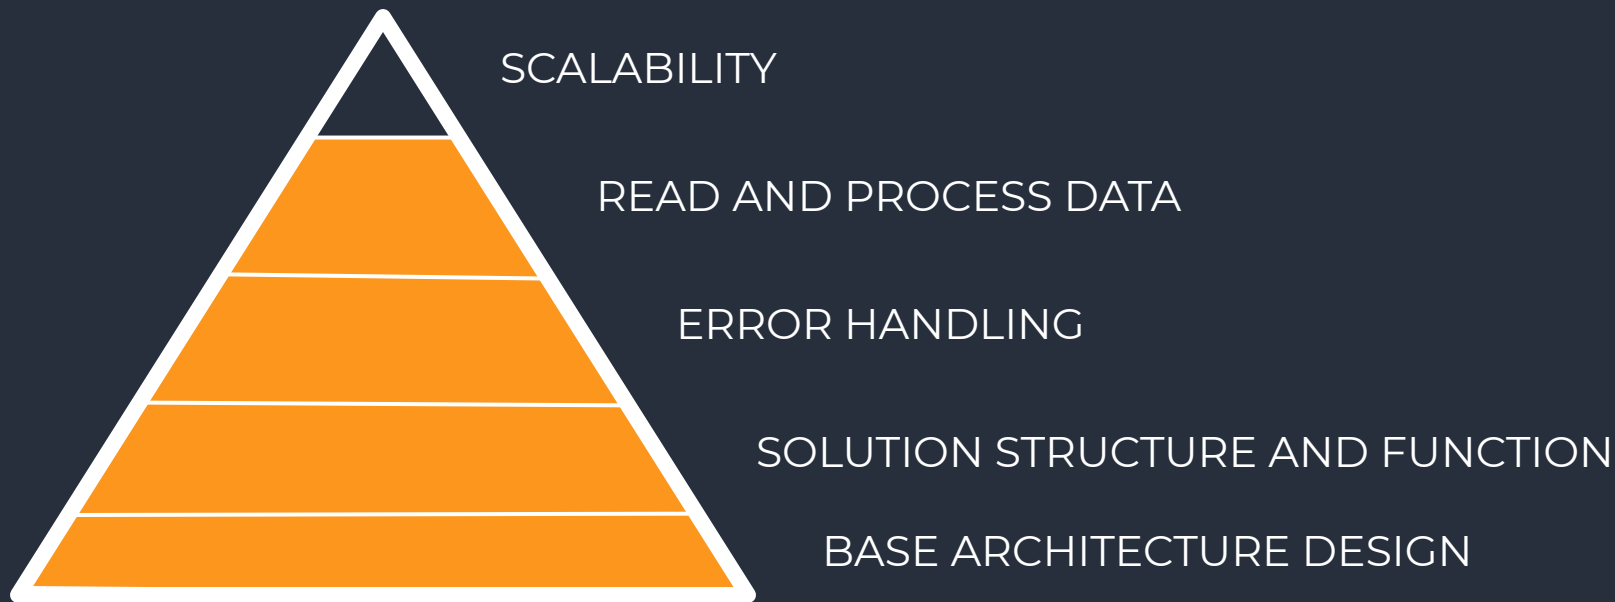
# ON-DEMAND REPROCESSING



# RESULT

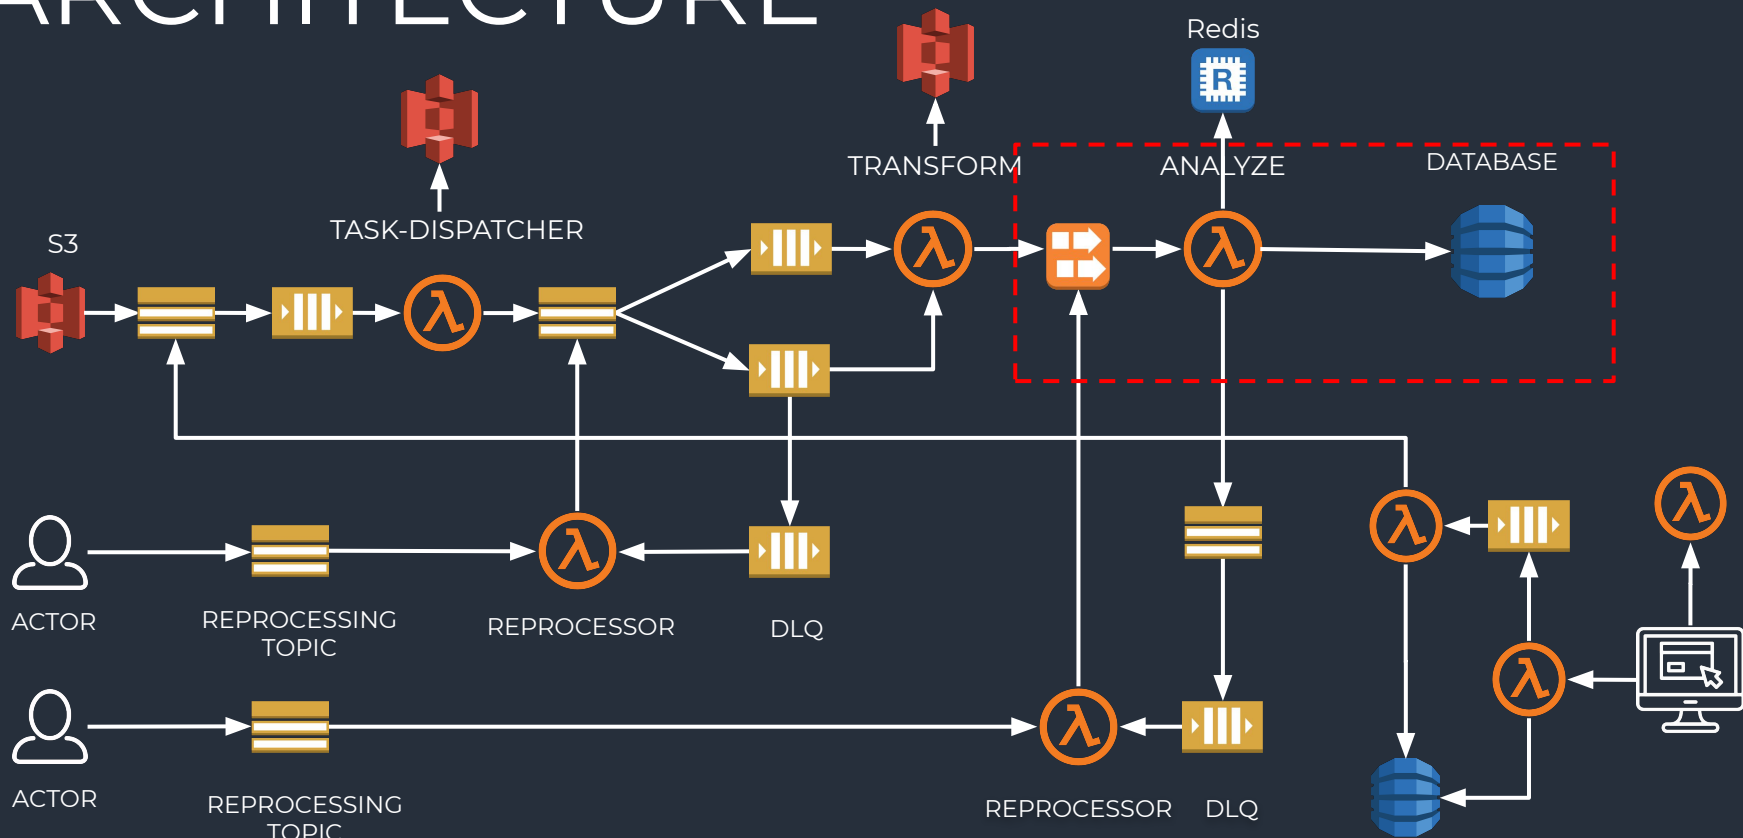


# PROGRESS

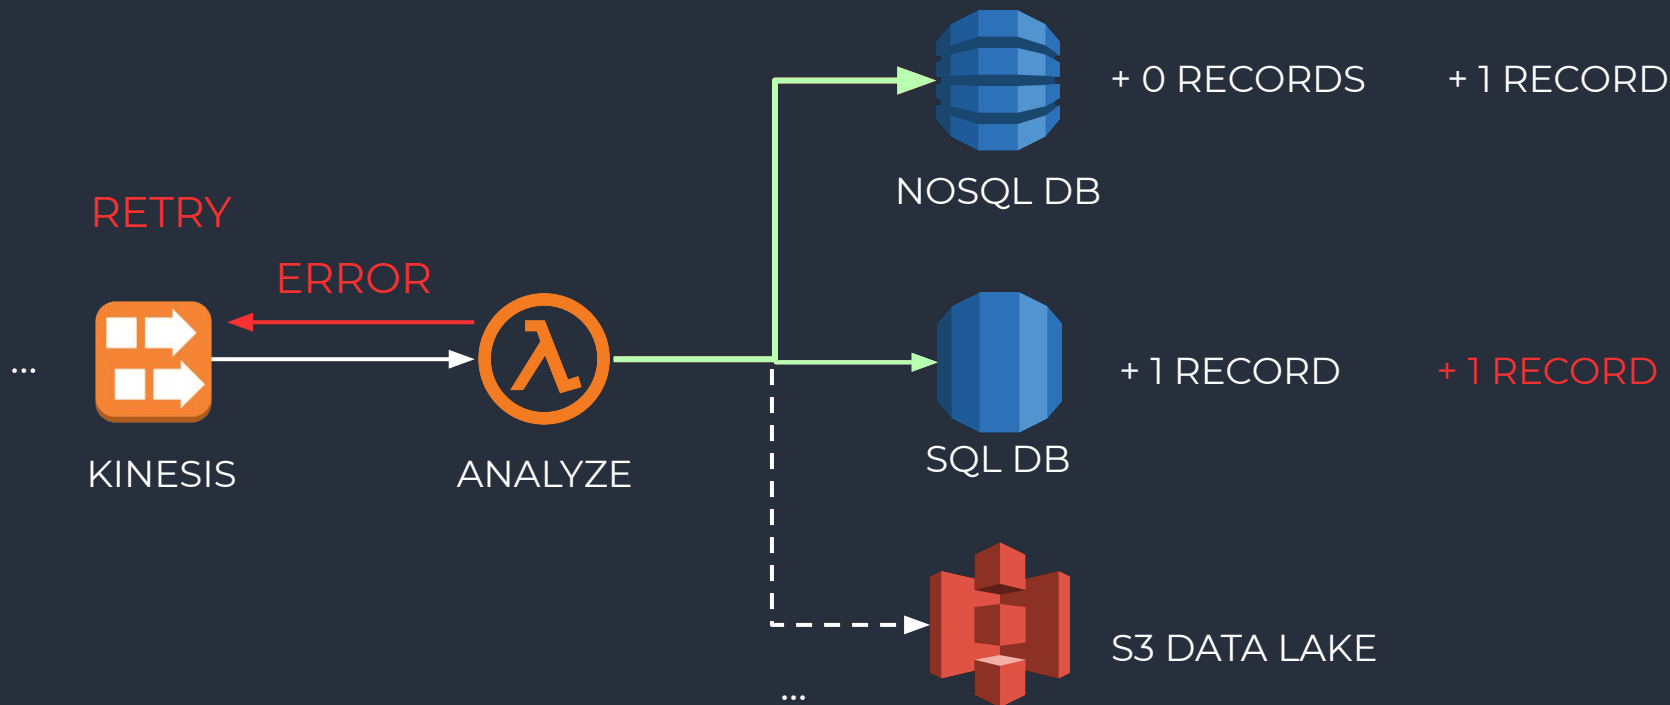




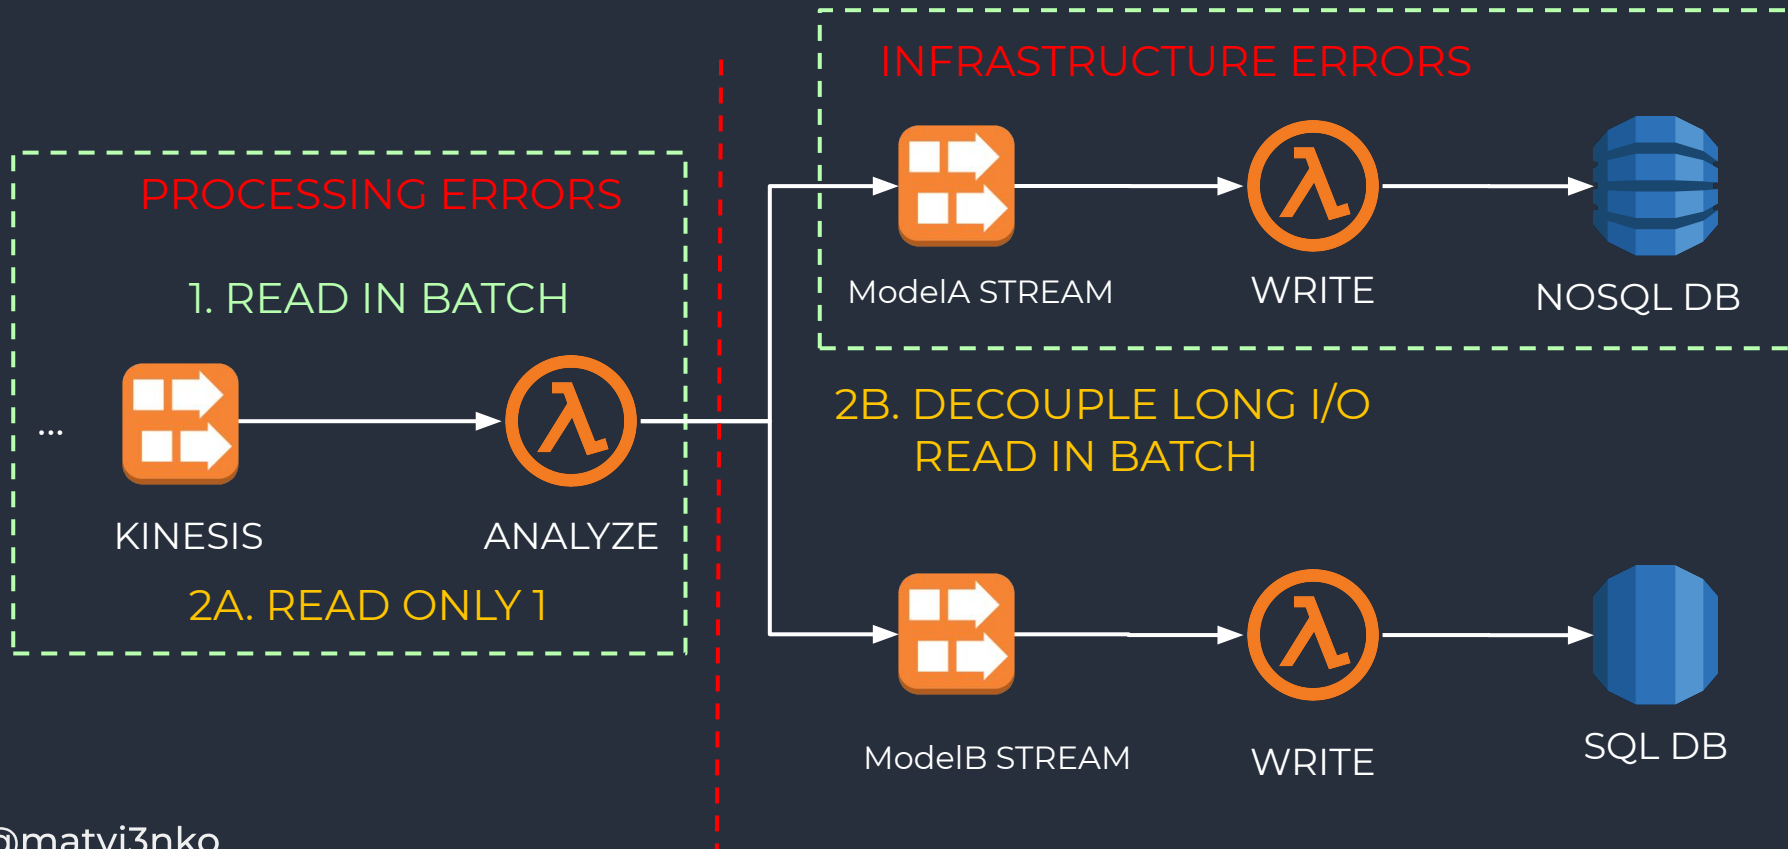
# ARCHITECTURE



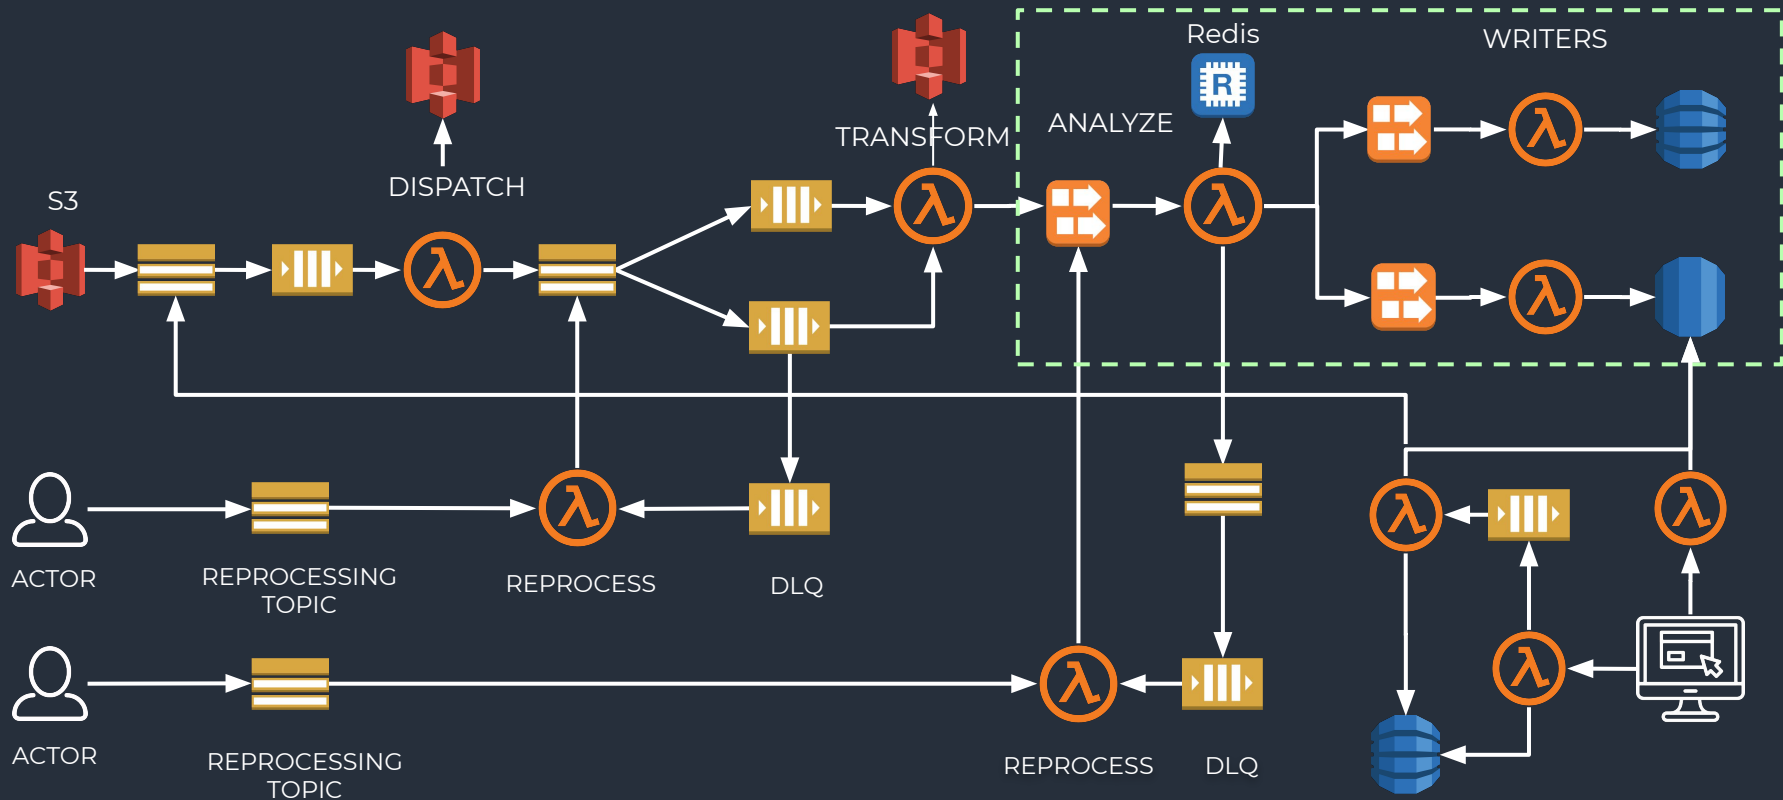
# WRITING TO A DATABASES



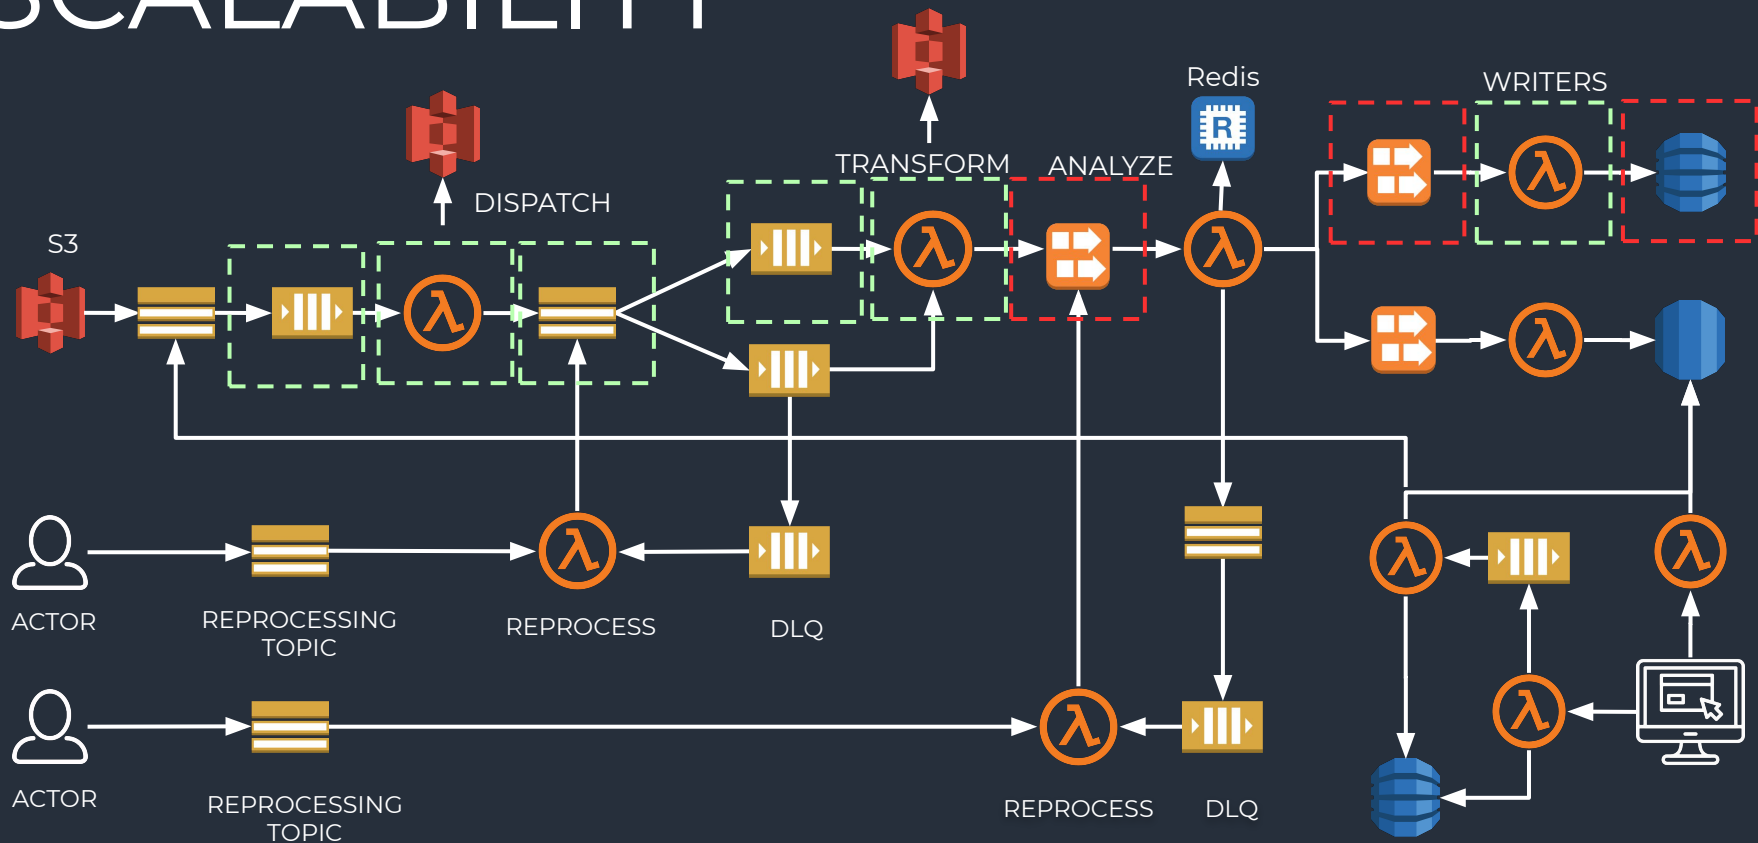
# DECOUPLING



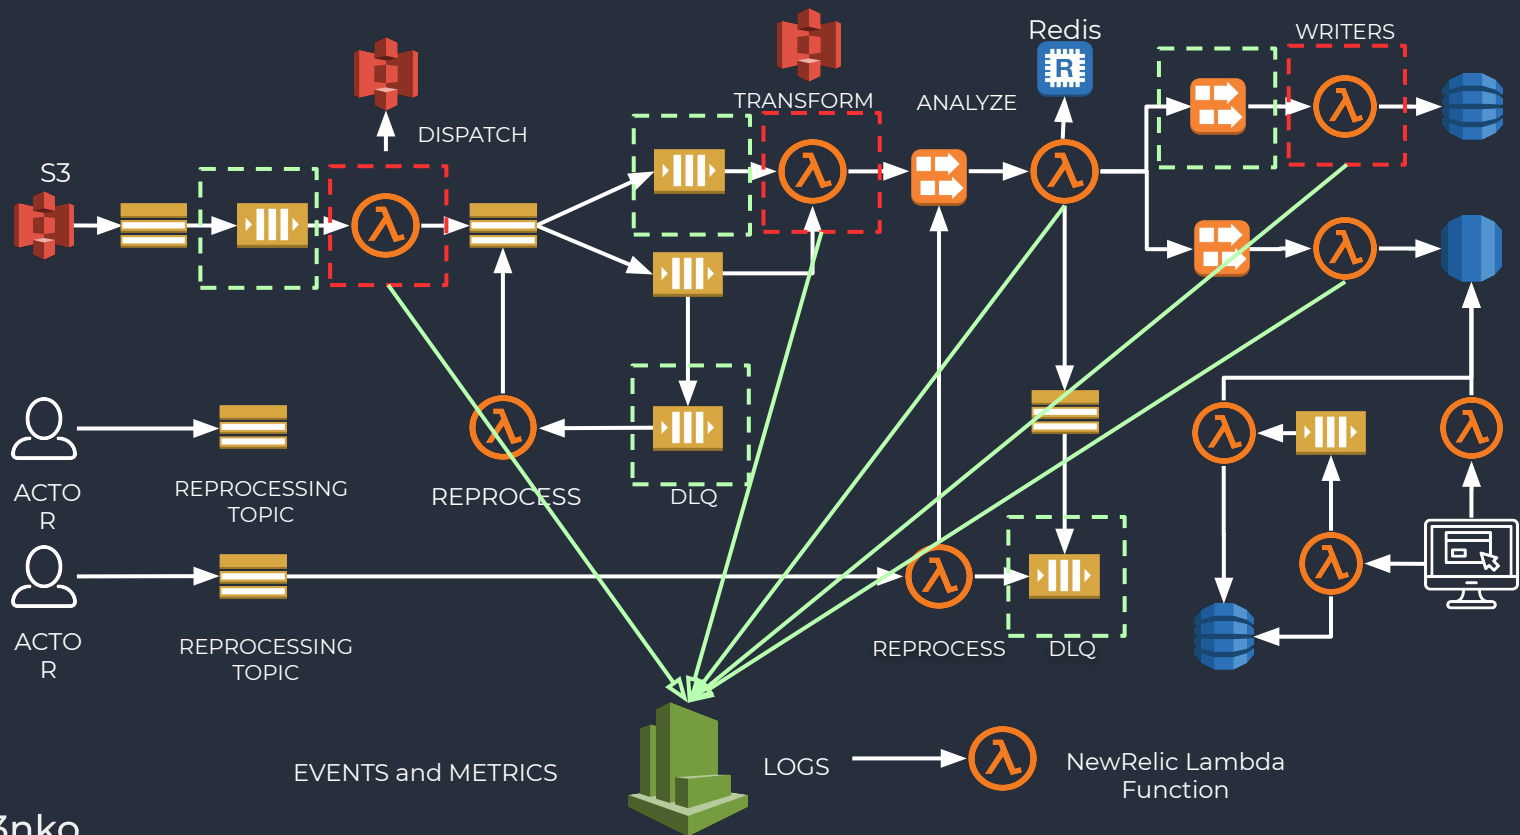
# DECOUPLED WRITERS



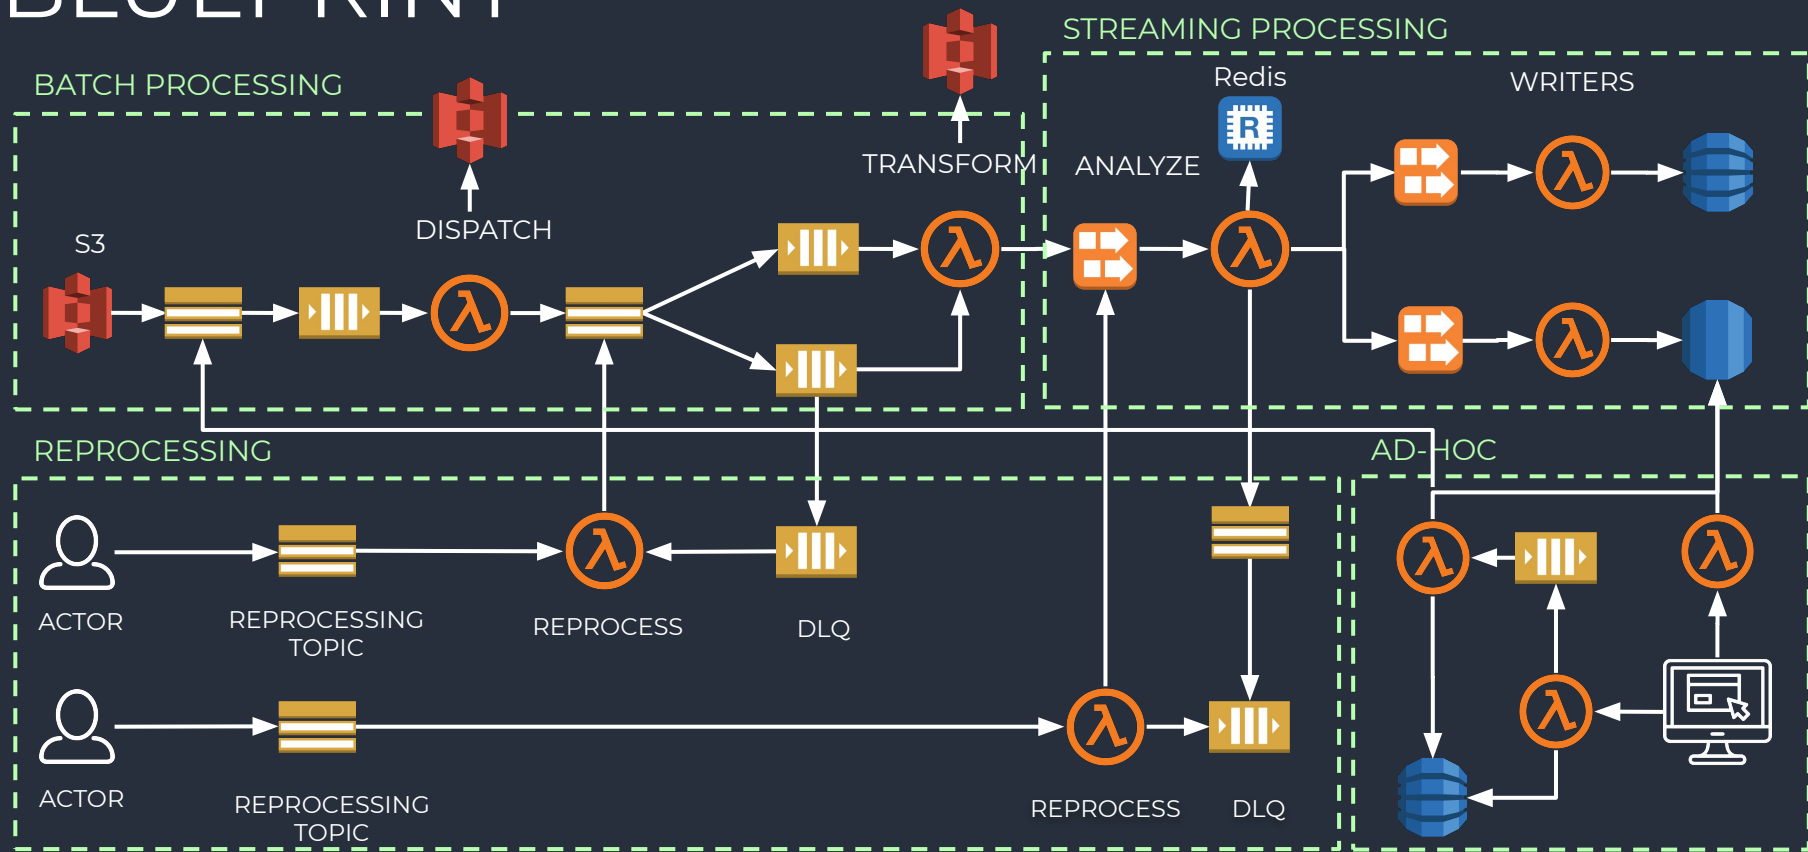
# SCALABILITY




# TROUBLESHOOTING



# BLUEPRINT



# PROGRESS



SCALABILITY

READ AND PROCESS DATA

ERROR HANDLING

SOLUTION STRUCTURE AND FUNCTION

BASE ARCHITECTURE DESIGN



# SUMMARY

1. Highly flexible architecture for changes in a result.
2. Massive scalable up to 5K, 7K, **10K lambda functions in parallel**
3. Terabytes of data
4. Serverless vs EC2: 60% – 5 times cheaper
5. Developer role in architecture

# THANKS!



**Nikolay Matvienko**

**matvi3nko@gmail.com**

**Twitter.com/matvi3nko**

**github.com/matvi3nko**